

# FinalProject

September 28, 2025

## 1 Modeling Shark Attack Occurrence: A Supervised Learning Approach

### Introduction/Background

Shark attacks are extremely rare events with great consequences for the victims involved. These events have become widely popularized in modern culture and media because of their violent and novel nature. That said, public exposure to these events widely inflated relative to their actual occurrence. Individual risk, as reported by the ISAF is 1 in 11.5 million, but this can vary based on the profile of the individual (e.g. individual's number of visits to the beach, activity while in the water, etc.).

With this in mind, shark attacks are still events that occur with risk that can be modeled. And where risk can be modeled, individuals can assess and make judgements for themselves on how to navigate through the world. This paper will explore various factors associated with shark attack occurrences, and will attempt to quantify the risk involved so that individuals can make judgements for themselves on how to interpret that risk.

The main sources of data for these events is The International Shark Attack File (ISAF). The ISAF is a catalog of all these events and is a widely trusted database on shark attack occurrences globally, curated by researchers at the University of Florida (special acknowledgement and thanks to Dr. Gavin Naylor and Joseph Miguez for this data). To help decrease model complexity and to provide a manageable project scope, only a subset of this data was pulled: specifically, unprovoked attack occurrences on the Atlantic Coast of Florida. The possibility of heterogeneity in pulling a specific dataset such as this, rather than global data is a limitation that will be discussed further in the conclusion of this paper.

Data Citations (also listed in references in bottom of doc):

Florida Museum. (n.d.). International Shark Attack File. Florida Museum of Natural History. Retrieved September 25, 2025, from <https://www.floridamuseum.ufl.edu/shark-attacks/>

University of West Florida, Haas Center. (n.d.). Tourism indicators. University of West Florida. Retrieved September 25, 2025, from <https://uwf.edu/centers/haas-center/explore-the-economy/tourism-indicators/>

### Problem Statement

The main problem that is being solved by this project is: Given temporal and weather data, can a shark attack occurrence be reliably predicted? This is not to say that the aim of this project is to model individual risk of attack, rather the general occurrence of an attack given the features at

hand. This is an important distinction that is necessary to understand when viewing the data and results.

## Hypothesis

The two general questions tested in this project are:

1. Can temporal features (i.e. time of day, month of year, and lunar phases) help provide indication of a shark attack occurrence?
2. Can weather features (i.e. barometric pressure, temperature, and the change of these features over time) provide indication of shark attack occurrence?

Regarding temporal features, it is widely known shark attacks occur more often on a seasonal basis. These events happen more often in warmer months due to more beachgoers and shark migratory patterns. Additionally, it is known that more attacks occur during day hours due to the fact that more people visit the beach during the day rather than night. Furthermore, it is also thought that the lunar phases have an effect on shark attack occurrence, due to general lunar illumination. These features and their effect on shark attack occurrence will be tested via the models later mentioned.

Regarding weather features, the interaction that is of interest for this project is barometric pressure, temperature, and the delta and rate of change between these variables over time. As all aquatic life are sensitive to dramatic shifts in their environment, these models will test whether or not these shifts influence shark behavior in a way that has an effect on shark attack occurrence. The causal nature of this affect is not being measured (i.e. whether sharks become more aggressive or whether they flee an area when there are dramatic shifts in these features), rather the overall correlation between these whether events and shark attack occurrences themselves.

## 2 1. Data Import and Cleansing

In terms of size, the data is 95 columns wide and has 736 rows, some of which have incomplete data.

First importing all the necessary libraries

```
[1]: import scipy as sp
import scipy.stats as stats
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import copy
import statsmodels.formula.api as smf
import statsmodels.api as sm
import re
import math
from datetime import datetime

#import ephem
```

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.metrics import classification_report, roc_auc_score,
    →average_precision_score, roc_curve, confusion_matrix
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline

```

Installing openpyxl as the CSV reader was giving me trouble

```
[2]: !pip install openpyxl
```

```

Collecting openpyxl
  Downloading openpyxl-3.1.3-py2.py3-none-any.whl.metadata (2.5 kB)
Collecting et-xmlfile (from openpyxl)
  Downloading et_xmlfile-1.1.0-py3-none-any.whl.metadata (1.8 kB)
Downloading openpyxl-3.1.3-py2.py3-none-any.whl (251 kB)
      251.3/251.3 kB
72.9 MB/s eta 0:00:00
Downloading et_xmlfile-1.1.0-py3-none-any.whl (4.7 kB)
Installing collected packages: et-xmlfile, openpyxl
Successfully installed et-xmlfile-1.1.0 openpyxl-3.1.3

```

```

[notice] A new release of pip is
available: 23.3.2 -> 24.0
[notice] To update, run:
pip install --upgrade pip

```

One of the main challenges in this project was data sanitization. The ISAF is a near-century old project, with data going back from the 1500s- present. Data drift is an inevitable threat with a project as long-standing as this, and sanitization is a normal part of the modeling process. Why and how values are inputted have changed over time, and the data reflects this. Multiple cleansing and parsing functions were used to create a sanitized output.

The first of these were functions created to parse through the ‘time-of-attack’ feature, which is one of the most critical to this analysis. As the data varied between written phrases such as “Forenoon” and “Afternoon”, the usage of AM/PM, military time, and general ranges such as “2:00- 4:00”, a variety of functions were created to parse this information. Additionally, functions were created to “bucket-ize” this time of attack feature and create new features based off of this. This includes the creation of a general time-of-day feature that shows whether the attack occurred in morning, afternoon, evening, or night, which was useful for visualizations.

Additional features were also created from the ‘date\_at\_location’ feature to help create more modeling value. The first of these was the creation of general months, and seasons, based on the date. The second was the creation of features related to the lunar phases. Three functions were created to calculate the general lunar phase (new, first quarter, full, last quarter), the general lunar illumination fraction (expressed as a fraction for modeling purposes), and general lunar illumination percentage (expressed as a percentage for visualization purposes).

Another thing to note is that unnecessary columns were dropped (such as body parts involved in the attack) to reduce noise while modeling and to curtail the dataframe. Additionally, all provoked attacks were dropped from the data, as the scope of this project pertains to solely unprovoked instances. As provoked instances vary widely, the risk is that they would have skewed the results of the analysis.

One of the byproducts of this project is a cleaned output of data, as well as a data cleansing pipeline that can both be used for further projects such as this.

**In terms of size, the data is 95 columns wide and has 736 rows, some of which have incomplete data.**

Reading in the data below

```
[3]: df = pd.read_excel('data/Atlantic_Coast_Florida_Metadata.xlsx',
    ↪engine='openpyxl')
df_2 = pd.read_excel('data/Air_Traffic_data.xlsx', engine='openpyxl')
```

Inspecting the data to see features available and data quality

```
[4]: df.head()
```

```
[4]:
```

	authenticity	attack_classification	\
0	Confirmed, shark involvement also confirmed	Unprovoked attack	
1	Confirmed, shark involvement also confirmed	Unprovoked attack	
2	Confirmed, shark involvement also confirmed	Unprovoked attack	
3	Confirmed, shark involvement also confirmed	Unprovoked attack	
4	Confirmed, shark involvement also confirmed	Unprovoked attack	

	outcome	continent	country	\
0	Non-fatal	North America	USA	
1	Fatal Attack, body or parts thereof recovered,...	North America	USA	
2	Non-fatal	North America	USA	
3	Non-fatal	North America	USA	
4	Non-fatal	North America	USA	

	state	county	\
0	Florida	Palm Beach County	
1	Florida	Brevard County	
2	Florida	Duval County	
3	Florida	St. Johns County	
4	Florida	St. Lucie County	

	0-all attacks.locality	date_at_location	Year	\
0	NaN	1931-09-21	1931	
1	Indiatlantic Beach, just across the bridge for...	1934-06-20	1934	
2	Mayport	1944-05-31	1944	
3	Crescent Beach, in front of Sea Haven Condos	1952-06-02	1952	
4	Fort Pierce, South Beach	1957-02-05	1957	

	...	bone_exposed	body_cavity_exposed	\
0	...	No	Data insufficient for judgement	
1	...	Data insufficient for judgement		No
2	...	Data insufficient for judgement		No
3	...	No		No
4	...	Data insufficient for judgement		No

	appendage_loss_to_shark	appendage_loss_to_surgery	trunk_severed	\
0	No	No	No	
1	No	No	No	
2	No	No	No	
3	No	No	No	
4	No	No	No	

	swallowed_whole	skeletonized	wounds_other_than_above	\
0	No	No	Data insufficient for judgement	
1	No	No	Data insufficient for judgement	
2	No	No	No	
3	No	No	Data insufficient for judgement	
4	No	No	No	

	provocative_acts	\
0	None known/unprovoked attack	
1	None known/unprovoked attack	
2	None known/unprovoked attack	
3	None known/unprovoked attack	
4	None known/unprovoked attack	

	activity_addenda
0	NaN
1	NaN
2	NaN
3	NaN
4	Floating on back, no flotation device reported

[5 rows x 95 columns]

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 736 entries, 0 to 735
```

```
Data columns (total 95 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	authenticity	736 non-null	object
1	attack_classification	736 non-null	object

2	outcome	736 non-null	object
3	continent	736 non-null	object
4	country	736 non-null	object
5	state	736 non-null	object
6	county	736 non-null	object
7	0-all attacks.locality	728 non-null	object
8	date_at_location	736 non-null	datetime64[ns]
9	Year	736 non-null	int64
10	victim_activity	736 non-null	object
11	abdomen_stomach	736 non-null	object
12	hands	736 non-null	object
13	arms	736 non-null	object
14	feet	736 non-null	object
15	calf_knee	736 non-null	object
16	head	736 non-null	object
17	chest	736 non-null	object
18	shoulder	736 non-null	object
19	thigh	736 non-null	object
20	buttocks	736 non-null	object
21	genitals	736 non-null	object
22	waist	736 non-null	object
23	fingers_toes	736 non-null	object
24	fingers	736 non-null	object
25	toes	736 non-null	object
26	back	736 non-null	object
27	injury_other_than_above	736 non-null	object
28	total_depth	735 non-null	object
29	depth_of_attack	736 non-null	object
30	relative_depth	622 non-null	object
31	habitat	736 non-null	object
32	turbidity	514 non-null	object
33	visibility	527 non-null	object
34	sunlight_conditions	692 non-null	object
35	gen_weather	563 non-null	object
36	air_temp	555 non-null	object
37	water_temp	736 non-null	object
38	water_clarity	736 non-null	object
39	sea_surf_cond	626 non-null	object
40	locality_Addenda	138 non-null	object
41	time_of_attack	733 non-null	object
42	closest_phase	736 non-null	object
43	waxing_waning	736 non-null	object
44	genus	736 non-null	object
45	species	736 non-null	object
46	common_name	736 non-null	object
47	prior_behavior	671 non-null	object
48	initial_behavior	736 non-null	object
49	subsequent_behavior	736 non-null	object

50	final_behavior	736 non-null	object
51	form_of_initial_strike	506 non-null	object
52	strike_direction	736 non-null	object
53	passes	736 non-null	object
54	discrete_strikes	736 non-null	object
55	discrete_bites	621 non-null	object
56	object_of_attack	731 non-null	object
57	shark_Addenda	358 non-null	object
58	no_of_sharks	701 non-null	object
59	scuba	736 non-null	object
60	face_mask	736 non-null	object
61	swim_fins	736 non-null	object
62	weapons	736 non-null	object
63	life_jacket	736 non-null	object
64	dive_bag	736 non-null	object
65	wetsuit_drysuit	736 non-null	object
66	swim_suit	736 non-null	object
67	non_swim_clothing	736 non-null	object
68	hard_hat_suit	736 non-null	object
69	gear_clothes_other_than_above	736 non-null	object
70	degree_type_of_clothing	573 non-null	object
71	primary_color	736 non-null	object
72	secondary_color	736 non-null	object
73	tertiary_color	736 non-null	object
74	gear_clothes_Addenda	333 non-null	object
75	surfboard_bottom_color	336 non-null	object
76	surfboard_size	321 non-null	object
77	surfboard_number_fins	33 non-null	object
78	clothing_pattern	578 non-null	object
79	special_features	580 non-null	object
80	scrapes_abrasions	736 non-null	object
81	discontinuous_tooth_marks	736 non-null	object
82	nonsevere_laceration	736 non-null	object
83	severe_lacerations	736 non-null	object
84	significant_tissue_loss	736 non-null	object
85	bone_exposed	736 non-null	object
86	body_cavity_exposed	736 non-null	object
87	appendage_loss_to_shark	736 non-null	object
88	appendage_loss_to_surgery	736 non-null	object
89	trunk_severed	736 non-null	object
90	swallowed_whole	736 non-null	object
91	skeletonized	736 non-null	object
92	wounds_other_than_above	736 non-null	object
93	provocative_acts	708 non-null	object
94	activity_addenda	231 non-null	object

dtypes: datetime64[ns](1), int64(1), object(93)

memory usage: 546.4+ KB

[ ]:

After doing an inspection of the data, it looks like there are a lot of features that I don't necessarily need for this project. All the data regarding the type of bite, body parts affected, and other miscellaneous information I removed so as to not confound any further analysis I do.

**In terms of size, the data is 95 columns wide and has 736 rows, some of which have incomplete data.**

```
[6]: df_clean = df[['attack_classification', 'outcome', 'county',  
    ↳ 'date_at_location', 'Year',  
    ↳ 'time_of_attack', 'victim_activity', 'sunlight_conditions',  
    ↳ 'gen_weather', 'air_temp', 'water_temp',  
    ↳ 'closest_phase', 'waxing_waning', 'common_name',  
    ↳ 'provocative_acts', 'activity_addenda']].copy()  
  
print(df_clean['time_of_attack'].unique())  
print(df_clean['outcome'].unique())  
print(df_clean['attack_classification'].unique())
```

```
['Data insufficient for judgement' '09:00'  
'data insufficient for judgement' '10:30' '1530' '11:00' '1700' '15:30'  
'15:00' '12:00' '11:30' 'late afternoon/evening' '1330' '14:30' '13:30'  
'07:00' '1500' '16:00' '16:30' 'afternoon' '14:00' '16:56' '08:55'  
'16:40' '13:50' '10:00' '16:06' '13:00' '11:15' 'morning' '13:15' '17:30'  
'14:45' '14:15' '19:30' '18:00' '09:45' '08:00' '12:30' 'forenoon'  
'17:00' '14:10' '09:15' '19:00' 'night' '18:45' '11:45' '18:30' '12:15'  
'18:20' '16:50' '11:50' '08:30' '16:35' '04:00' '10:55' '17:45' '07:45'  
'18:12' '14:25' '19:35' '09:25' '14:54' '09:10' '13:20' '14:20' '08:45'  
'19:07' '15:50' '07:30' '18:15' '11:06' '13:25' '18:05' '13:45' '11:20'  
'11:36' '10:25' '09:30' '13:43' '17:15' '18:02' '09:35' '11:10' '17:59'  
'20:15' '15:45' '18:40' 'Afternoon' '11:57' '14:55' '17:02' '14:50'  
'17:42' '12:48' '11:00-14:00' '12:46' '10:50' '13:06' '13:14' '16:15'  
'16:20' '14:40' '10:15' '06:45' '10:40' '14:47' '12:05' '15:55' '13:40'  
'19:15' '12:20' '13:23' '13:53' '11:55' '10:07' '14:37' '17:11' '16:14'  
'06:00' '17:25' '11:12' '17:05' '13:55' '14:01' '16:45' nan '14:21'  
'13:21' '12:37' '18:43' '16:12' '11:05' '13:42' '08:10' '18:50' '12:03'  
'15:10' '11:02' '11:40' '10:45' '10:23' '12:45' '13:18' '13:17' '9:50'  
'14:12' '10:48' '08:15' '19:23' '08:14' '11:48' '17:18' '14:27' '14:46'  
'12:53' '15:40' '12:38' '09:50' '14:59' '18:19' '15:38' '15:36' '10:32'  
'15:28' '1000' '2 PM' '5:30 pm' '1930' '1030' '0758' '1036' '1630'  
'12:25' '17:03' '17:20' '10:47' '15:20' '9:00' '7:00' '16:48' '15:57'  
'16:38' '12:23' '10AM' '06:30' '18:36' '13,20' '10:41' '11:14' '15:18'  
'11:25' '16:16' '15:05' '14:06' '9:20' '7:44' '12:11' '9:30' '18:04'  
'15:52']
```

```
['Non-fatal']
```

```
'Fatal Attack, body or parts thereof recovered, death considered direct result
```



```

of shark-inflicted wounds'
'Assumed fatal attack, body not recovered, no personal gear recovered'
'Assumed fatal attack, body not recovered, personal gear recovered'
'Fatal, body or parts thereof recovered, not known whether death was a direct
result of shark-inflicted wounds']
['Unprovoked attack' 'Provoked attack' 'No assignment can be made'
'Air/sea disaster: unprovoked attack']

```

As you can see below, I am cleansing the “Outcome” feature below. I am creating a binary and ternary feature that describes whether the outcome was fatal or non-fatal based on the text entered in this column. This will be useful if I decide to do logistic modeling on the Outcome

Also, I am deleting all rows where the attack classification is a provoked attack. This culls 64 rows from the overall dataset, but this is a necessary action due to the fact that provoked attacks could happen for too many numerous factors. This will protect our data integrity

```

[7]: #Data Cleansing

df_clean = df_clean[df_clean['attack_classification'] != 'Provoked attack'].
    ↪copy()

#deleting all rows that are provoked attacks

df_clean['outcome'] = df_clean['outcome'].replace({
'Non-fatal': 'non-fatal',
'Fatal Attack, body or parts thereof recovered, death considered direct result
    ↪of shark-inflicted wounds': 'fatal',
'Assumed fatal attack, body not recovered, no personal gear recovered':
    ↪'assumed_fatal',
'Assumed fatal attack, body not recovered, personal gear recovered':
    ↪'assumed_fatal',
'Fatal, body or parts thereof recovered, not known whether death was a direct
    ↪result of shark-inflicted wounds': 'fatal'
})

for col in ['attack_classification', 'outcome', 'county', 'victim_activity',
    ↪'sunlight_conditions', 'gen_weather',
            'closest_phase', 'waxing_waning', 'common_name', 'provocative_acts',
    ↪'activity_addenda']:
    df_clean[col] = df_clean[col].astype(str).str.strip().str.lower()

df_clean['outcome_binary'] = df_clean['outcome'].map({'fatal': 1,
    ↪'assumed_fatal': 1, 'non-fatal': 0})

df_clean['outcome_ternary'] = df_clean['outcome'].map({'fatal': 1,
    ↪'assumed_fatal': 2, 'non-fatal': 0})

```

```
[8]: print(df_clean['outcome'].unique())
```

```
['non-fatal' 'fatal' 'assumed_fatal']
```

Now, I am getting into the temporal data cleansing below. I created a few functions that go through the `time_of_attack` feature and normalize it so it is consistent throughout. I really would like this feature to be used in modelling so I am spending a lot of time and effort cleaning it.

Additionally I am creating a function that bucketizes the data into general day phases (i.e. dawn, day, dusk, night)

```
[9]: #time_of_attack needs some heavy data cleansing here

def midpoint(h1, h2):
    if h1 is None or h2 is None:
        return None
    return (h1 + ((h2-h1)% 24)/2) % 24

def ampmparse(s: str):
    for fmt in ("%I%p", "%I %p", "%I:%M%p", "%I:%M %p"):
        try:
            t = datetime.strptime(s.strip().upper(), fmt)
            return t.hour + t.minute/60.0
        except ValueError:
            continue
    return None

def hhmmparse(hhmm: str):
    try:
        hh, mm = hhmm.split(":")
        hh, mm = int(hh), int(mm)
        if 0 <= hh <= 23 and 0 <= mm <= 59:
            return hh + mm/60.0
    except Exception:
        pass
    return None

def normalize_time(val):
    if val is None or (isinstance(val, float) and math.isnan(val)):
        return np.nan
    s = str(val).strip()

    sl = s.lower()
    if sl == "data insufficient for judgement":
        return np.nan

    s = s.replace(",", ":")
```

```

#for the single instance of midpoint range 11:00 - 14:00
if "-" in s:
    try:
        left, right = s.split("-")
        h1, h2 = hhmmparse(left.strip()), hhmmparse(right.strip())
        return midpoint(h1, h2)
    except Exception:
        return np.nan

#ampm normalization
ampm_val = ampparse(s)
if ampm_val is not None:
    return ampm_val

#military time normalization
if re.fullmatch(r"\d{3,4}", s):
    hh, mm = int(s[:-2]), int(s[-2:])
    if 0<= hh <= 23 and 0<= mm <= 59:
        return hh + mm/60.0

# HH:MM and H:MM
if re.fullmatch(r"\d{1,2}:\d{2}", s):
    return hhmmparse(s)

# HH
if re.fullmatch(r"\d{1,2}", s):
    hh = int(s)
    if 0<= hh <= 23:
        return float(hh)

#General phrases
phrase_map = {"forenoon": 10.0, "morning": 9.0, "afternoon": 15.0, "late_
→afternoon/evening" : 18.0, "night": 22.0}
if s1 in phrase_map:
    return phrase_map[s1]

return np.nan

def time_buck(h):

    if pd.isna(h):
        return "unknown"
    if 5 <= h < 9:
        return "dawn"
    if 9 <= h < 17:
        return "day"

```

```

if 17<= h < 21:
    return "dusk"
return "night"

```

Creating my new time of attack features from the functions above

```

[10]: #data cleansing contd.

df_clean["time_of_attack_hour"] = df_clean["time_of_attack"].
    ↪apply(normalize_time)

df_clean["time_unknown_flag"] = df_clean["time_of_attack_hour"].isna().
    ↪astype(int)

df_clean["time_bucket"] = df_clean["time_of_attack_hour"].apply(time_buck)

```

Now, just printing all the unique values to make sure that this worked correctly and to get an idea of what I am working with

```

[11]: print(df_clean['time_of_attack_hour'].unique())
print(df_clean['time_bucket'].unique())
print(df_clean['time_unknown_flag'].unique())

```

```

[      nan  9.      10.5      15.5      11.      17.
 15.      12.      11.5      18.      14.5      13.5
 7.      16.      16.5      14.      16.93333333  8.91666667
16.66666667 13.83333333 10.      16.1      13.      11.25
13.25      17.5      14.75      19.5      9.75      8.
12.5      14.16666667 19.      22.      18.75      11.75
18.5      12.25      14.25      18.33333333 16.83333333 11.83333333
 8.5      16.58333333 4.      10.91666667 17.75      7.75
18.2      14.41666667 9.41666667 14.9      9.16666667 13.33333333
14.33333333 8.75      19.11666667 15.83333333 7.5      18.25
11.1      13.41666667 18.08333333 13.75      11.33333333 11.6
10.41666667 9.5      13.71666667 17.25      18.03333333 9.58333333
11.16666667 17.98333333 20.25      15.75      18.66666667 14.91666667
17.03333333 14.83333333 17.7      12.8      12.76666667 10.83333333
13.1      13.23333333 16.25      16.33333333 14.66666667 10.25
 6.75      10.66666667 14.78333333 12.08333333 15.91666667 13.66666667
19.25      12.33333333 13.38333333 13.88333333 11.91666667 14.61666667
17.18333333 16.23333333 6.      17.41666667 11.2      17.08333333
13.91666667 14.01666667 16.75      14.35      13.35      12.61666667
18.71666667 16.2      11.08333333 13.7      8.16666667 12.05
15.16666667 11.03333333 11.66666667 10.75      10.38333333 12.75
13.3      13.28333333 9.83333333 14.2      10.8      8.25
19.38333333 18.83333333 8.23333333 11.8      17.3      14.45
14.76666667 12.88333333 15.66666667 12.63333333 14.98333333 18.31666667
15.63333333 15.6      15.46666667 10.6      12.41666667 17.05

```

```

17.33333333 10.78333333 15.33333333 16.8          15.95          16.63333333
12.38333333 6.5          18.6          10.68333333 11.23333333 15.3
16.26666667 15.08333333 14.1          9.33333333 7.73333333 12.18333333
18.06666667 15.86666667]
['unknown' 'day' 'dusk' 'dawn' 'night']
[1 0]

```

Now, I want to work on cleaning the date at location feature, deriving the month and general season from that data as well, so I can test my hypothesis on seasonality.

There is some odd air\_temp data that seems to be in Fahrenheit, so I am masking that here as well. This may not be used, but I will cleanse anyways.

```

[12]: df_clean['date_at_location'] = pd.to_datetime(df_clean['date_at_location'],
    ↪errors = 'coerce')
df_clean['month'] = df_clean['date_at_location'].dt.month
df_clean['season'] = df_clean['month'].map({
    12: 'winter', 1: 'winter', 2: 'winter',
    3: 'spring', 4: 'spring', 5: 'spring',
    6: 'summer', 7: 'summer', 8: 'summer',
    9: 'fall', 10: 'fall', 11: 'fall'})

for col in ['air_temp', 'water_temp']:
    df_clean[col] = pd.to_numeric(df_clean[col], errors = 'coerce')

faren_mask = df_clean['air_temp'] > 60
df_clean.loc[faren_mask, 'air_temp'] = (df_clean.loc[faren_mask, 'air_temp'] -
    ↪32) *(5/9)

hr = df_clean['time_of_attack_hour']

df_clean['hour_sin'] = np.sin(2*np.pi*hr/24)
df_clean['hour_cos'] = np.cos(2*np.pi*hr/24)

```

In the cell below I create functions that derive the moon phase and lunar illumination from the date. This should be more reliable than the given lunar features in the data, and will help me test my hypothesis on lunar phase impacting the occurrence of shark attacks

As you can see above, I end up with 4 new features-

lunar\_phase: which bucketizes all the occurrences into 4 equal bins (new, last\_quarter, first\_quarter, and full)

lunar\_illumination\_phase: the 0-1 value of lunar phase

lunar\_illumination\_frac: 0-1 value that indicates the lunar illumination (0-1 easier for modeling)

lunar\_illumination\_pct: 0-100 percentage value that indicates the lunar illumination (0-100 easier for visualization)

```

[13]: def moon_phase(date):
        if pd.isna(date):
            return None
        last_known_new_moon = pd.Timestamp("2000-01-06")
        syn_month = 29.5305887
        days = (pd.to_datetime(date) - last_known_new_moon).days
        return (days % syn_month) / syn_month

df_clean["lunar_illumination_phase"] = df_clean["date_at_location"].
    ↪ apply(moon_phase)

def illuminationfrac(phase):
    if phase is None:
        return None
    return 0.5 * (1 - np.cos(2 * np.pi * phase))

df_clean["lunar_illumination_frac"] = df_clean["lunar_illumination_phase"].
    ↪ apply(illuminationfrac)
df_clean["lunar_illumination_pct"] = df_clean["lunar_illumination_frac"] * 100.0

def moon_phase_norm(val):
    if val is None:
        return "unknown"
    if val < 0.125 or val >= 0.875:
        return "new"
    elif val < 0.375:
        return "first_quarter"
    elif val < 0.625:
        return "full"
    else:
        return "last_quarter"

df_clean["lunar_phase"] = df_clean["lunar_illumination_phase"].
    ↪ apply(moon_phase_norm)

print(df_clean['lunar_phase'].value_counts(dropna = False))
print(df_clean['lunar_illumination_pct'].value_counts(dropna = False))
print(df_clean['lunar_illumination_frac'].value_counts(dropna = False))
print(df_clean['lunar_illumination_phase'].value_counts(dropna = False))

```

new	183
full	166
last_quarter	163
first_quarter	159

```

Name: lunar_phase, dtype: int64
95.828611    4
0.170482     4
12.160051    4
1.573899     4
0.422984     3
..
99.692238    1
72.243376    1
97.194988    1
11.465200    1
74.325882    1
Name: lunar_illumination_pct, Length: 599, dtype: int64
0.121601     4
0.015739     4
0.958286     4
0.001705     4
0.004230     3
..
0.920522     1
0.058864     1
0.500471     1
0.011702     1
0.408360     1
Name: lunar_illumination_frac, Length: 600, dtype: int64
0.040039     4
0.565472     4
0.979283     3
0.013147     3
0.113381     2
..
0.280508     1
0.749577     1
0.756782     1
0.999479     1
0.264208     1
Name: lunar_illumination_phase, Length: 604, dtype: int64

```

Now, I think I am cleansing done for now, so I would like to test and visualize some of the counts and values of the data in the below cells

```

[14]: print(df_clean['month'].unique())
      #print(df_clean['date_at_location'].unique())
      print(df_clean['air_temp'].unique())
      print(df_clean['water_temp'].unique())

```

```
[ 9  6  5  2  4  3  7  8  1 10 11 12]
[      nan 24.      26.      29.      28.      30.
 18.      27.      31.      21.      32.      33.
 20.      23.      22.      26.7      28.5      31.1
 25.6      35.      23.8      16.      29.4      32.22
 37.      28.8      25.5      32.2      27.7      33.3
 30.3      36.      28.9      30.5      25.      22.22
 23.89      27.2      19.      24.44      22.77      26.667
 28.3      21.7      31.66666667 31.7      30.6      27.8
 21.4      34.4      33.8      26.1      26.6      32.78
 22.7      29.44      23.9      31.11      ]
[  nan 20.      25.      22.      26.      21.      24.      28.      18.      23.
 29.      27.      17.      31.      32.      30.      33.      19.4      19.      25.6
 29.4      21.6      25.5      26.5      15.6      26.7      23.9      24.44      24.4      20.5
 23.8      22.2      27.8      27.6      21.1      20.6      27.2      26.1      25.56      26.67
 23.3      28.3      21.67      22.7      23.88      27.222      28.9      26.6      27.7      27.78
 27.22      28.33 ]
```

```
[15]: print(df_clean['outcome'].value_counts(dropna = False))
print(df_clean[['outcome', 'outcome_binary', 'outcome_ternary']].head(10))
print(df_clean['time_bucket'].value_counts(dropna = False))
df_clean[['air_temp', 'water_temp']].describe()
```

```
non-fatal      662
fatal           6
assumed_fatal   3
Name: outcome, dtype: int64
   outcome  outcome_binary  outcome_ternary
0  non-fatal              0                0
1    fatal              1                1
2  non-fatal              0                0
3  non-fatal              0                0
4  non-fatal              0                0
5  non-fatal              0                0
8  non-fatal              0                0
9  non-fatal              0                0
10 non-fatal              0                0
12 non-fatal              0                0
day      444
unknown   97
dusk      95
dawn       33
night       2
Name: time_bucket, dtype: int64
```

```
[15]:      air_temp  water_temp
count  214.000000  373.000000
```



mean	28.201139	25.713812
std	3.888885	3.163131
min	16.000000	15.600000
25%	26.000000	24.000000
50%	29.000000	26.000000
75%	31.000000	28.000000
max	37.000000	33.000000

Printing out info and data head below of my newly cleansed data. A couple things that stick out to me is that the hourly time of attack features, as well as the temperature features have a lot of null values. I may have to either omit those rows, or fill the data synthetically at a later point

```
[16]: df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 671 entries, 0 to 735
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   attack_classification                 671 non-null    object
1   outcome                             671 non-null    object
2   county                              671 non-null    object
3   date_at_location                     671 non-null    datetime64[ns]
4   Year                                671 non-null    int64
5   time_of_attack                       668 non-null    object
6   victim_activity                     671 non-null    object
7   sunlight_conditions                 671 non-null    object
8   gen_weather                         671 non-null    object
9   air_temp                           214 non-null    float64
10  water_temp                          373 non-null    float64
11  closest_phase                       671 non-null    object
12  waxing_waning                      671 non-null    object
13  common_name                        671 non-null    object
14  provocative_acts                   671 non-null    object
15  activity_addenda                   671 non-null    object
16  outcome_binary                     671 non-null    int64
17  outcome_ternary                    671 non-null    int64
18  time_of_attack_hour                 574 non-null    float64
19  time_unknown_flag                   671 non-null    int64
20  time_bucket                         671 non-null    object
21  month                              671 non-null    int64
22  season                             671 non-null    object
23  hour_sin                           574 non-null    float64
24  hour_cos                           574 non-null    float64
25  lunar_illumination_phase            671 non-null    float64
26  lunar_illumination_frac             671 non-null    float64
27  lunar_illumination_pct              671 non-null    float64
28  lunar_phase                         671 non-null    object
```

```
dtypes: datetime64[ns](1), float64(8), int64(5), object(15)
memory usage: 157.3+ KB
```

```
[17]: df_clean.head()
```

```
[17]:  attack_classification  outcome      county date_at_location  Year \
0      unprovoked attack  non-fatal  palm beach county    1931-09-21  1931
1      unprovoked attack    fatal    brevard county    1934-06-20  1934
2      unprovoked attack  non-fatal    duval county    1944-05-31  1944
3      unprovoked attack  non-fatal  st. johns county    1952-06-02  1952
4      unprovoked attack  non-fatal  st. lucie county    1957-02-05  1957

      time_of_attack \
0  Data insufficient for judgement
1                      09:00
2  data insufficient for judgement
3                      10:30
4                      1530

      victim_activity \
0                      swimming
1      standing still on bottom
2      data insufficient for judgement
3      standing still on bottom
4  floating, little or no motion (includes use of...

      sunlight_conditions      gen_weather  air_temp \
0                      daylight          clear      NaN
1                      daylight  data insufficient for judgement      NaN
2  data insufficient for judgement  data insufficient for judgement      NaN
3                      daylight  data insufficient for judgement      NaN
4                      daylight  data insufficient for judgement      NaN

...  time_unknown_flag  time_bucket  month  season  hour_sin  hour_cos \
0  ...                1    unknown    9    fall      NaN      NaN
1  ...                0        day    6    summer  0.707107 -0.707107
2  ...                1    unknown    5    spring      NaN      NaN
3  ...                0        day    6    summer  0.382683 -0.923880
4  ...                0        day    2    winter -0.793353 -0.608761

      lunar_illumination_phase  lunar_illumination_frac  lunar_illumination_pct \
0                      0.316535                      0.702989                      70.298907
1                      0.281316                      0.597749                      59.774941
2                      0.306293                      0.673187                      67.318656
3                      0.322267                      0.719313                      71.931296
4                      0.194463                      0.329044                      32.904399
```

```

    lunar_phase
0  first_quarter
1  first_quarter
2  first_quarter
3  first_quarter
4  first_quarter

[5 rows x 29 columns]

```

```
[18]: #df_clean.to_csv('df_clean_here', index = False)
```

## 3 2. EDA and Visualization on Temporal Data

Exploratory data analysis and the creation of visuals is a key step for a project such as this, because it allows one the ability to see what key features they are going to use in modelling and how these features might affect the model. Many visualizations were created to make sense of the data. The visualizations shared below are the most relevant to the modelling that was subsequently done.

### Temporal Visualizations

The below visualizations related to count of shark attack occurrences across different measures of time are quite simple in and of themselves, but they do show us a few different things to look out for in the subsequent modeling. The first is that, unsurprisingly, shark attacks for this dataset tend to skew towards daytime hours, specifically in the mid-day to afternoon. The second is that shark attacks for this dataset skew toward late summer, early fall months. The third is that shark attacks for this dataset may have a slight uptick around new moons. All this being said, these visualizations aren't enough to determine whether these imbalances are simply indicative of the dataset itself, or can be used to infer future occurrences. In other words, are these results statistically significant?

### 3.1 Correlation Heatmap

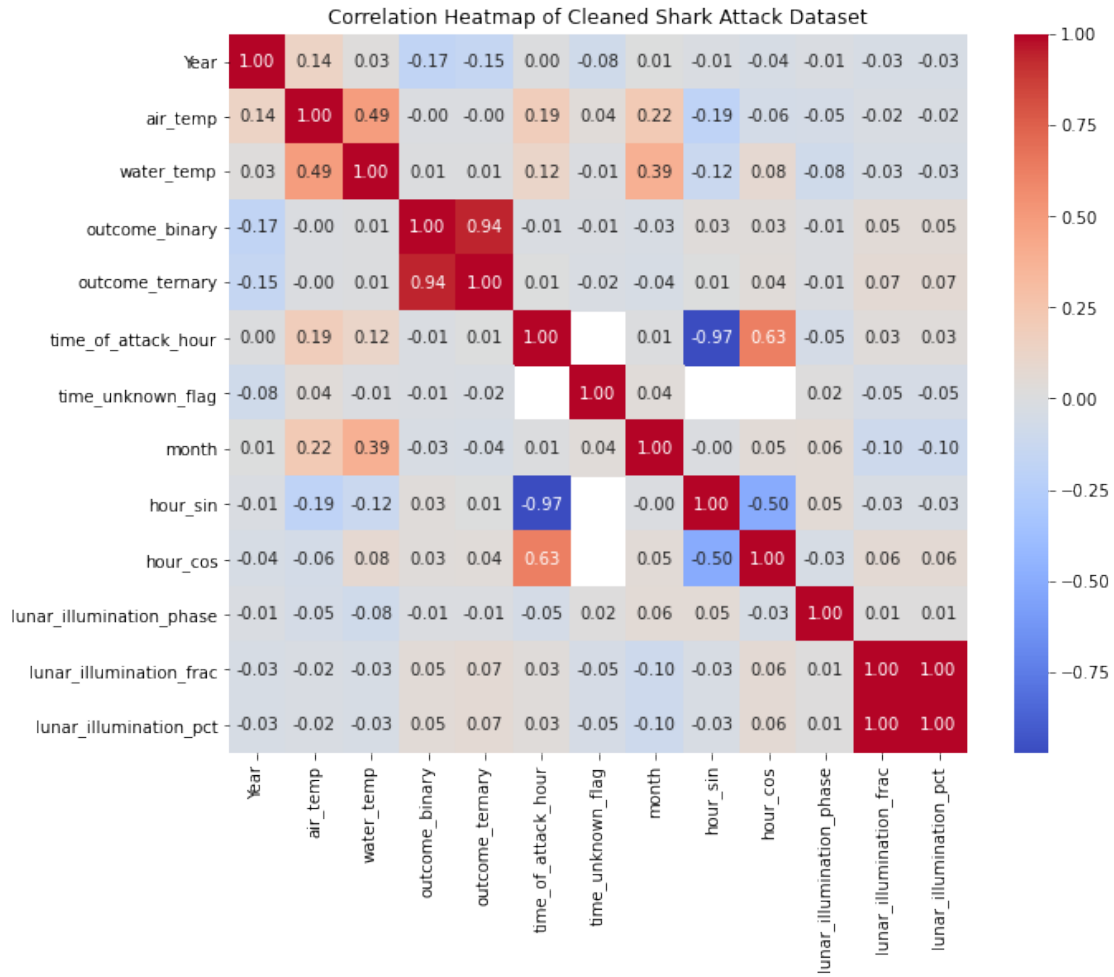
Looking at the features below, there are a few that have some correlation between them. The highest are between some of our feature-engineered values (e.g. hour-sin, hour-cos, and hour), but we can ignore those correlations for now because the engineered features are dependent on what they have been created from.

Some potentially interesting correlations below is water temp and month, as well as air temp and month. This makes sense intuitively as there is a correlation as the air temperature increases in times of the year.

At the moment, this correlation heatmap wont be extremely usefull in identifying model features, due to the fact that synthetic negatives are not created yet. This will be run again after the synthetic negatives are added in.

```
[19]: plt.figure(figsize=(10, 8))
sns.heatmap(df_clean.select_dtypes(include = ["number"]).corr(),
            annot = True, cmap = "coolwarm", fmt = ".2f", cbar = True)
```

```
plt.title("Correlation Heatmap of Cleaned Shark Attack Dataset")
plt.show()
```



## 3.2 Hourly Visualizations

The first thing I want to visualize is the shark attack occurrences vs hour of day.

As you can see below, a majority of shark attacks happen towards the daylight hours. This doesn't mean that sharks aren't willing to bite people at night-time though. This is more likely due to the fact that beachgoers attend the beach during daylight hours.

```
[20]: print(sns.__version__)

plt.figure()
sns.distplot(df_clean['time_of_attack_hour'], bins = 24, kde = False)
```

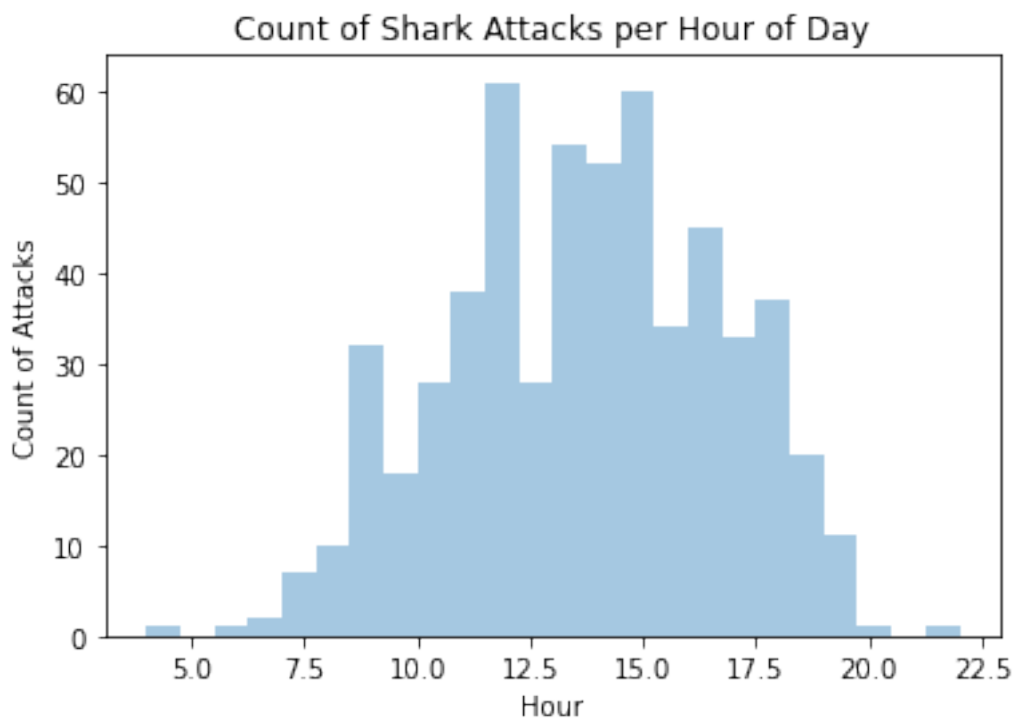
```

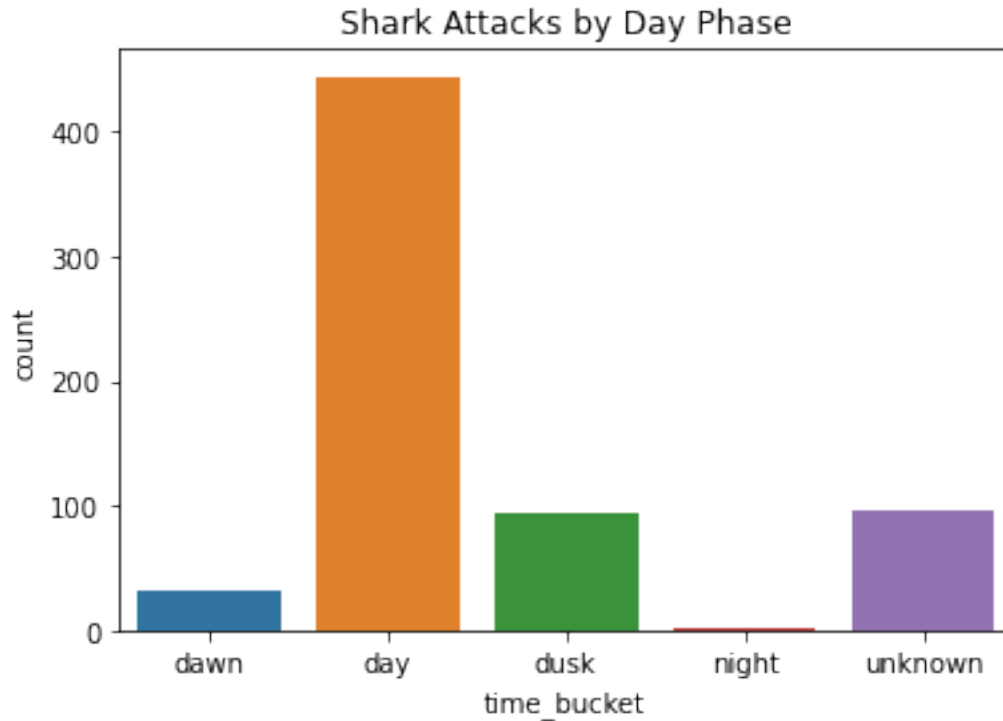
plt.xlabel("Hour")
plt.ylabel("Count of Attacks")
plt.title("Count of Shark Attacks per Hour of Day")
plt.show()

plt.figure()
sns.countplot(x = 'time_bucket', data = df_clean, order = ["dawn", "day", "dusk", "night", "unknown"])
plt.title("Shark Attacks by Day Phase")
plt.show()

```

0.10.1





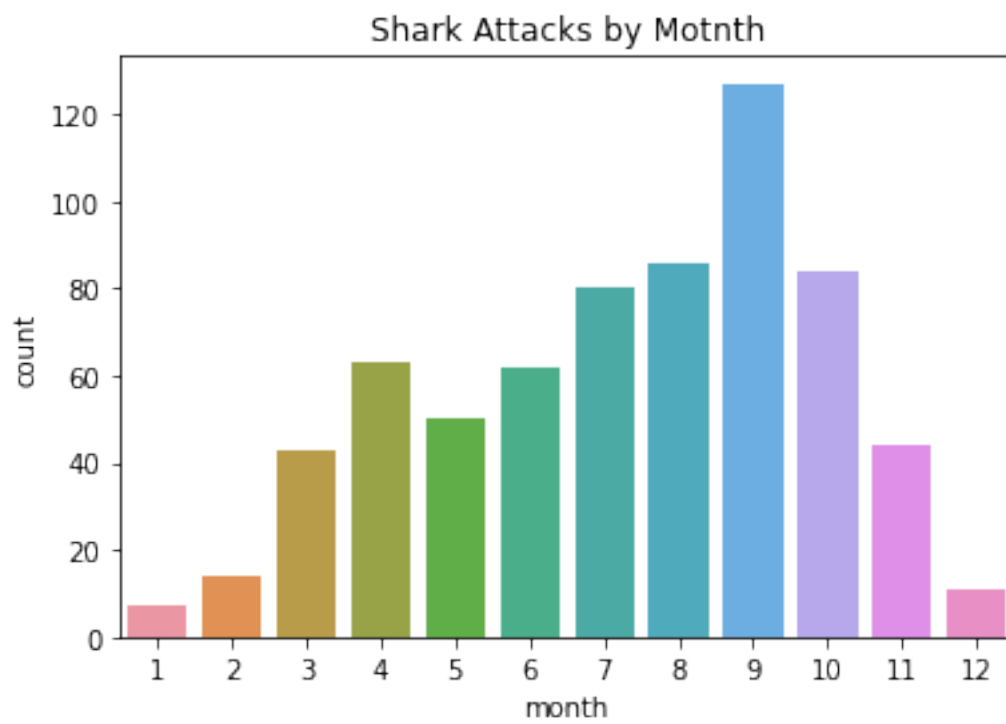
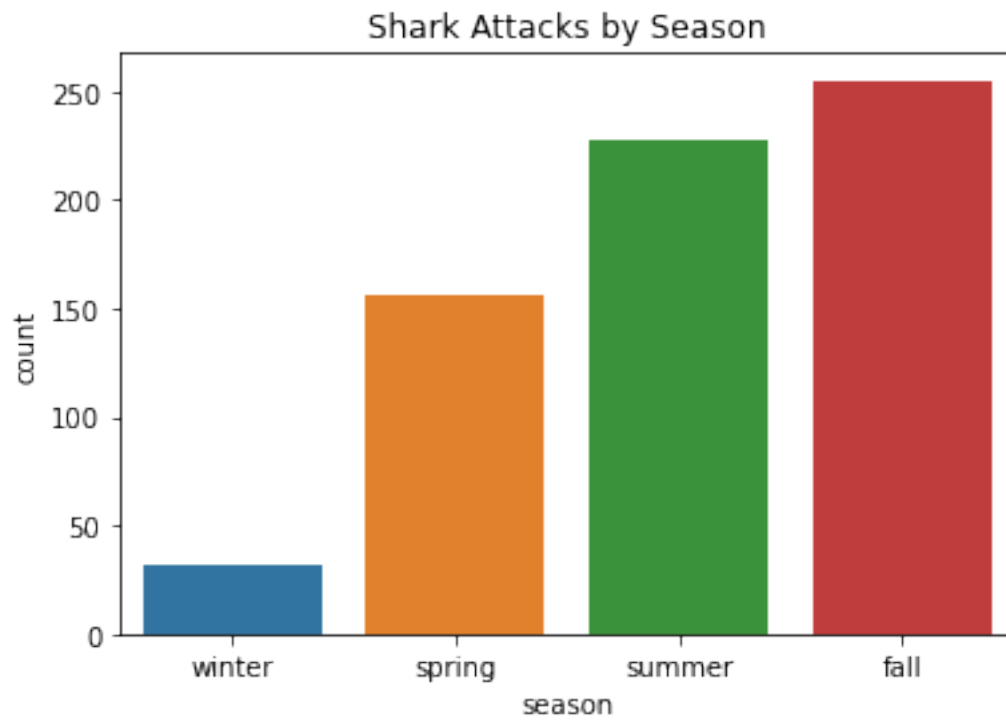
### 3.3 Monthly/Seasonal Visualizations

Again, visualizing some of the temporal data, but looking at it on a monthly/seasonal basis below.

As you can see, we have a large amount of attacks skewed toward the summer and fall months (months 6-11). Again, this could likely be due to the fact that beachgoers are more likely to visit the beach in the summer and fall months, when temperatures in Florida are more mild, but I would like to inspect this further.

```
[21]: plt.figure()
sns.countplot(x = 'season', data = df_clean, order = ["winter", "spring", "summer", "fall"])
plt.title("Shark Attacks by Season")
plt.show()

plt.figure()
sns.countplot(x = 'month', data = df_clean,)
plt.title("Shark Attacks by Month")
plt.show()
```



### 3.4 Lunar Phase/Illumination Visualizations

Now, I am looking at the interaction between shark attacks and lunar phase/illumination percentage.

Looking below, you can first see the shark attack occurrence plotted with the lunar illumination phase (mapped from 0-1, where 0 is a New moon, 0.5 is full, and 1 is a new moon again). This looks to show a slight skew toward the new moon, but this visualization is not great due to how the feature is set up. The feature is more useful for modeling than visualizing.

As you can see in the second graphic, I mapped the count of attacks by the lunar illumination percentage. This chart is interesting, but it can be deceiving, because it looks as though we have a very large skew toward new moons and full moons. While this seems to be the case, the reality is that if you were to also chart the lunar illumination by time, you would see a similar breakdown.

So, all that to say, I believe the final visual to be most telling. It shows the shark attack occurrences by bucketized versions of the lunar phases (first quarter, full, last quarter, and new moon). As you can see in the data, there seems to be a slight skew toward the new moon, but we cannot tell from this visualization alone whether that skew is significant or not. That is what our modeling will be for

```
[22]: plt.figure()
sns.distplot(df_clean['lunar_illumination_phase'].dropna(), bins = 20, kde = ☐
          ↪False )
plt.xlabel("Lunar Illumination with (0 = New Moon, 0.5 = Full, and 1 = New ☐
          ↪Moon)")
plt.ylabel("Count of Attacks")
plt.title("Count of Shark Attacks by Lunar Illumination Phase")
plt.show()

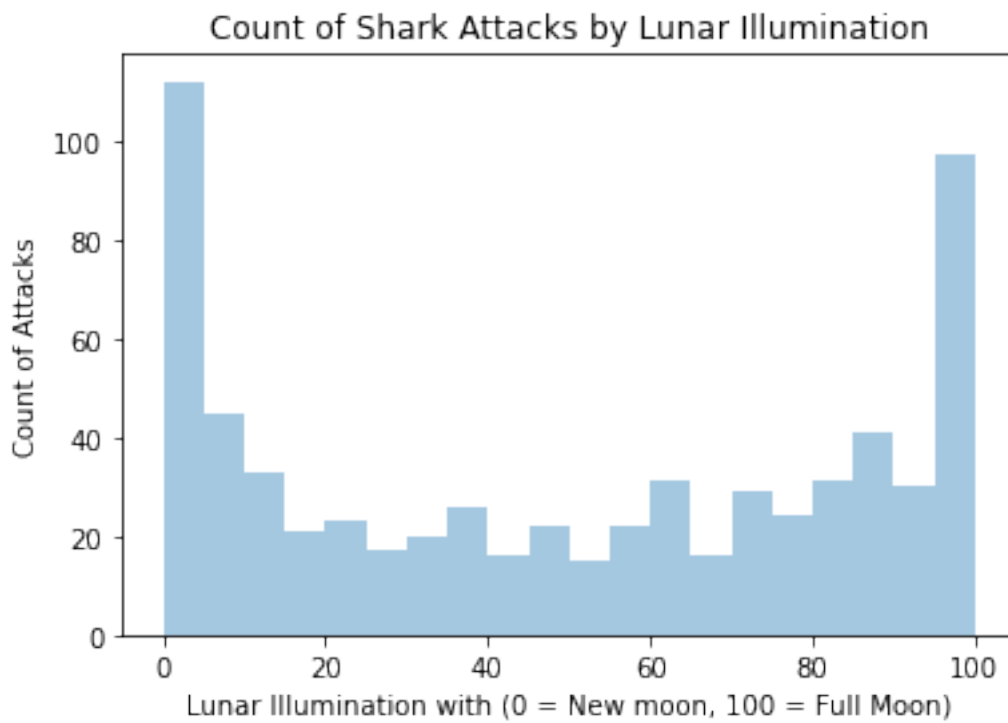
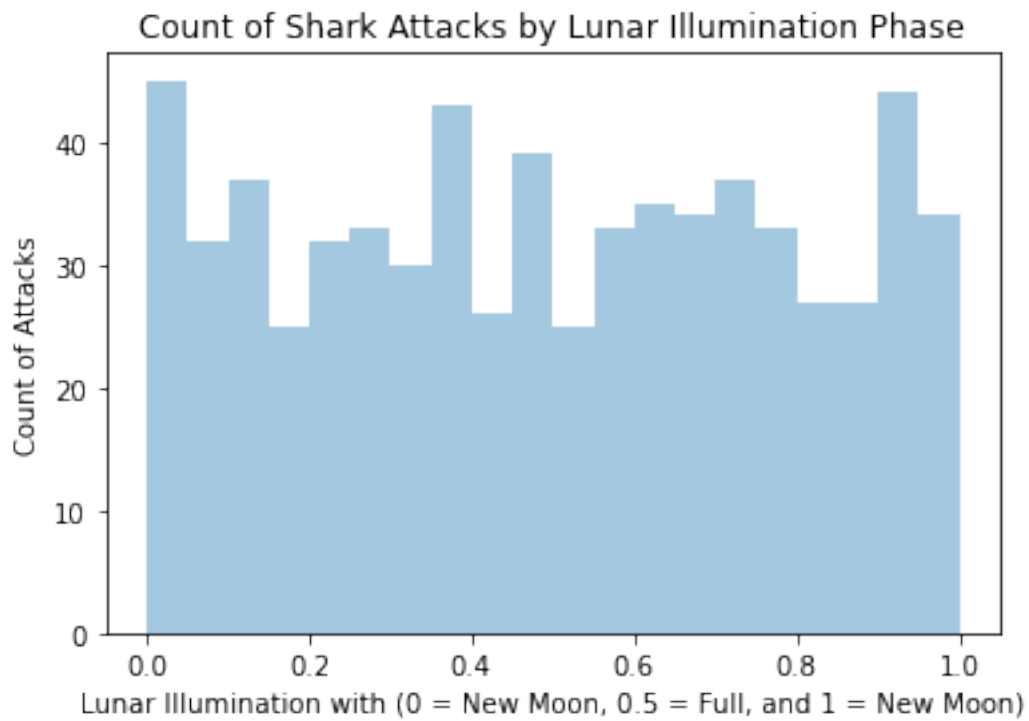
plt.figure()
sns.distplot(df_clean['lunar_illumination_pct'].dropna(), bins = 20, kde = ☐
          ↪False)
plt.xlabel("Lunar Illumination with (0 = New moon, 100 = Full Moon)")
plt.ylabel("Count of Attacks")
plt.title("Count of Shark Attacks by Lunar Illumination")
plt.show()

# plt.figure()
# sns.countplot(x = 'closest_phase', data = df_clean, order=☐
#               ↪sorted(df_clean['closest_phase'].dropna().unique()))
# plt.xticks(rotation = 45)
# plt.title("shark Attacks By Lunar Phase")

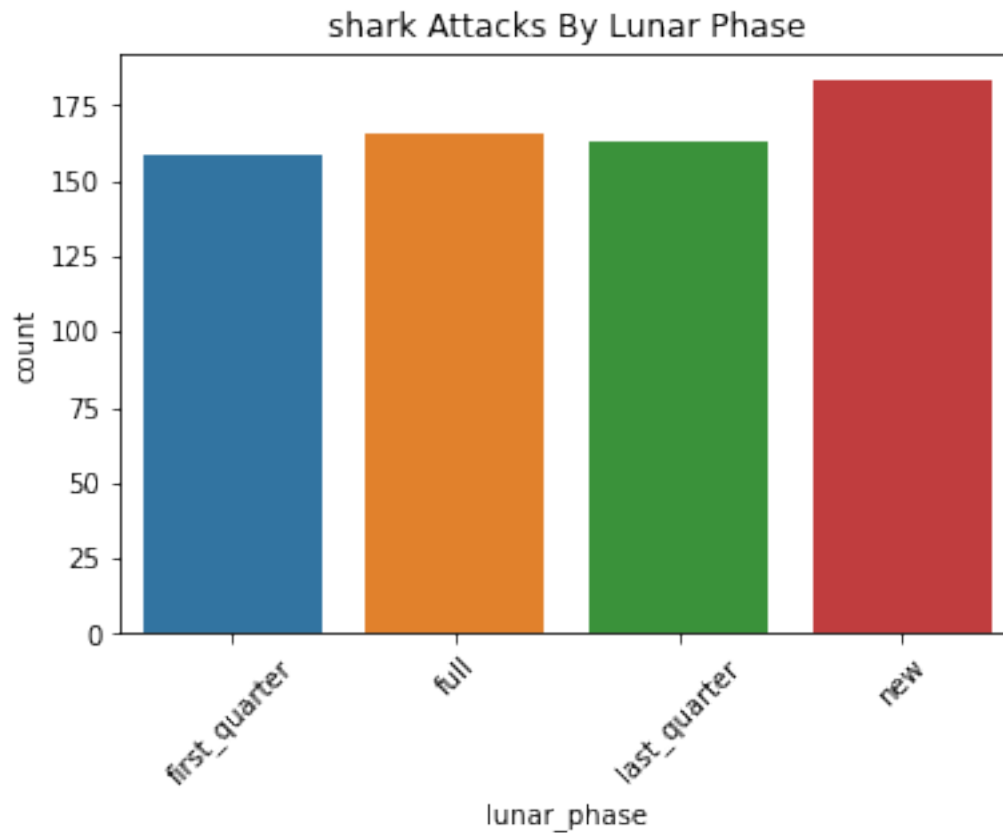
plt.figure()
sns.countplot(x = 'lunar_phase', data = df_clean, order=☐
          ↪sorted(df_clean['lunar_phase'].dropna().unique()))
plt.xticks(rotation = 45)
```



```
plt.title("shark Attacks By Lunar Phase")
```



```
[22]: Text(0.5, 1.0, 'shark Attacks By Lunar Phase')
```



### 3.5 Chi-Squared testing on Lunar Phases

Doing a Chi-square test to see if there is significance in moon phase distributions

As you can see by the results, the p value is quite high (.259). This means that the shark attacks across the lunar phases are not significant enough to tell me that it is more significant than an even spread.

Although my EDA showed that there was a slight skew toward the new moon, it is not significant enough to filter out the noise.

```
[23]: phase_counter = df_clean['lunar_phase'].value_counts()
      obsv = phase_counter.values

      expt = [obsv.sum()/len(obsv)]*len(obsv)

      chi2, p_value = stats.chisquare(f_obs = obsv, f_exp = expt)
```

```
print(obsv)
print(expt)
print(chi2)
print(p_value)
```

```
[183 166 163 159]
[167.75, 167.75, 167.75, 167.75]
1.9955290611028316
0.5733351830119464
```

### 3.6 Bonus Visualizations: Count of Attacks over time and shark attacks vs temperature

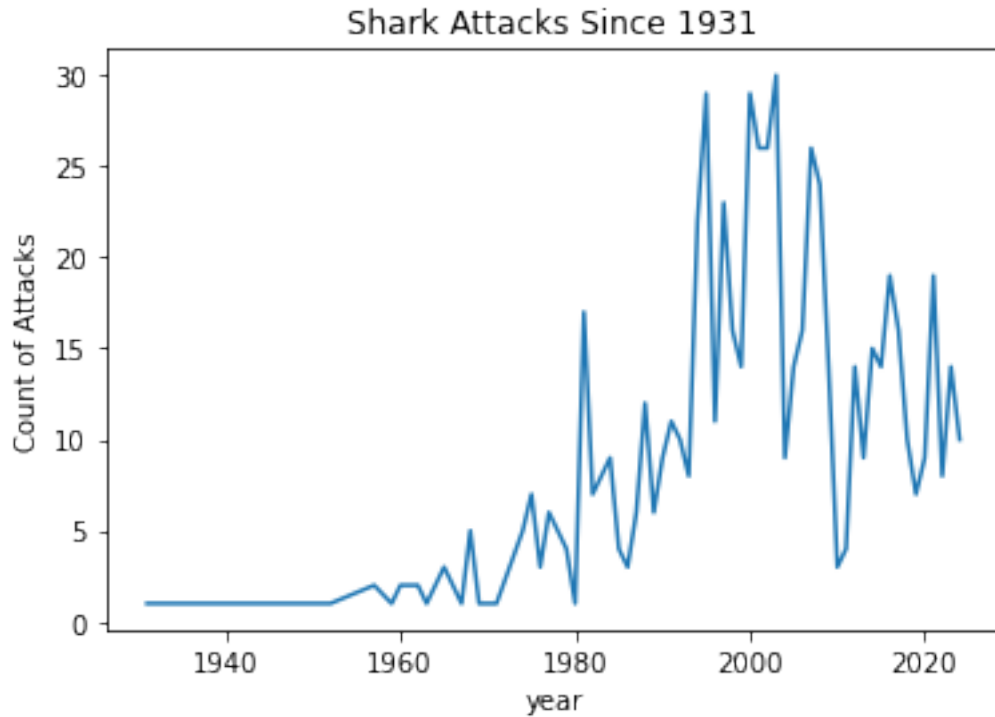
I created a few extra visualizations to see some other interactions that might be interesting to note.

As you can see, the count of shark attacks over time has been increasing. This is a well-known phenomenon and largely attributed to the fact that our global population has grown. More people = more bites. If we were to map this data to the population of Florida, I am sure we would find a similar trend.

Additionally, we have a count of water temperature variance and shark attack occurrence. This is an interesting thing to note and something we may come back to later. A skew toward a smaller range of temperatures, but this may just because the majority of Florida seasons involve these temperatures

```
[22]: plt.figure()
df_clean.groupby('Year').size().plot()
plt.xlabel("year")
plt.ylabel("Count of Attacks")
plt.title("Shark Attacks Since 1931")
```

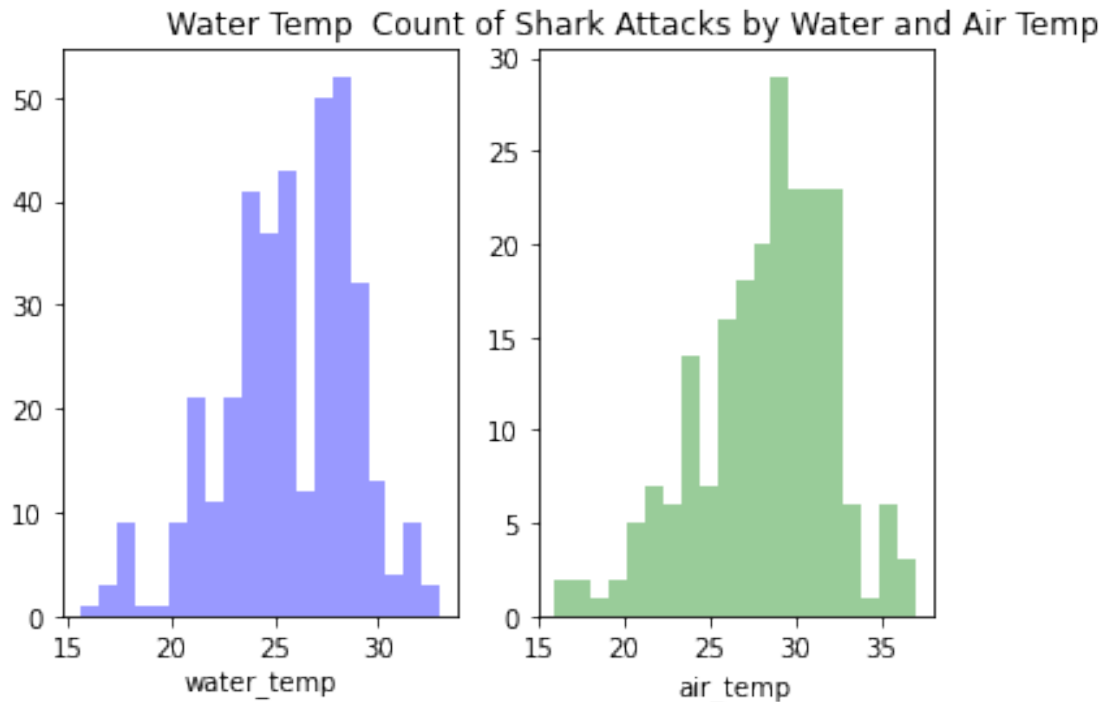
```
[22]: Text(0.5, 1.0, 'Shark Attacks Since 1931')
```



```
[23]: fig = plt.figure()
axes = fig.subplots(1,2)
sns.distplot(df_clean['water_temp'].dropna(), bins = 20, ax = axes[0], color = "blue", kde = False)
axes[0].set_title("Water Temp")
sns.distplot(df_clean['air_temp'].dropna(), bins = 20, ax = axes[1], color = "green", kde = False)
axes[1].set_title("Air Temp")
plt.title("Count of Shark Attacks by Water and Air Temp")
plt.show

#Maybe try overlaying month data below air and month temp
```

```
[23]: <function matplotlib.pyplot.show(*args, **kw)>
```



```
[24]: print(df_clean[['time_bucket', 'time_of_attack_hour', 'month',
↳ 'lunar_illumination_frac']].isna().sum())
```

```
time_bucket          0
time_of_attack_hour  97
month                0
lunar_illumination_frac  0
dtype: int64
```

### 3.7 2.5 Bonus: Modeling Shark attack Seasonality via Linear regression

Just out of curiosity, I wanted to create a linear regression model with month predicting shark attack. As you can see below, it is a very deficient model. A single predictor like this is not good enough, and it doesn't fit a linear regression model. I even sin/cos encoded the months to account for a more cyclical action, but this did little to help.

The below model is more of an example of why certain models don't fit certain datasets more than anything else. It shows that the initial thought to use logistic regression for this particular problem is the right one.

```
[25]: monthly_count = df_clean.groupby('month').size().reset_index(name =
↳ 'attack_count')

#monthly_count['month'] = df_clean['month'].copy()
```

```

monthly_count['sin_month'] = np.sin(2*np.pi*monthly_count['month'])
monthly_count['cos_month'] = np.cos(2*np.pi*monthly_count['month'])

X = monthly_count[['sin_month', 'cos_month']]
y = monthly_count['attack_count']

model_linear = LinearRegression()
model_linear.fit(X, y)

print(model_linear.coef_)
print(model_linear.intercept_)

monthly_count['pred'] = model_linear.predict(X)

plt.figure()
plt.plot(monthly_count['month'], monthly_count['attack_count'], 'o-', label_
↪='Actual')
plt.plot(monthly_count['month'], monthly_count['pred'], 'r--', label_
↪='Predicted')

plt.xlabel('Month')
plt.ylabel('Attack Counts by Month via Linear Regression')
plt.legend()
plt.show()

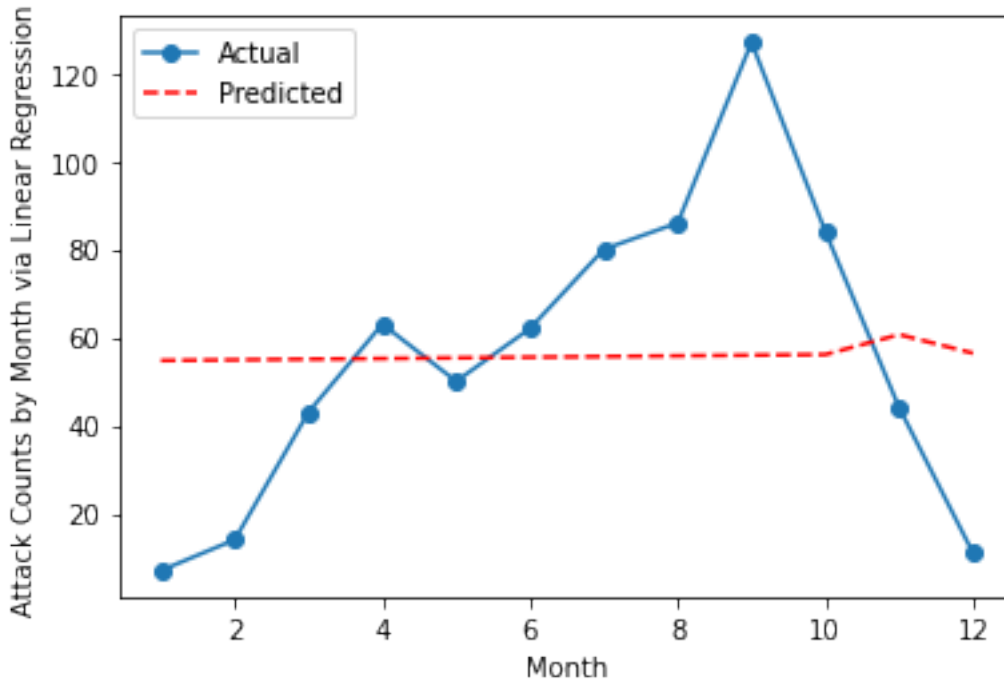
monthly_count.info()

```

```

[-6.17507605e+14  0.00000000e+00]
54.56793138810292

```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12 entries, 0 to 11
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   month            12 non-null    int64
1   attack_count     12 non-null    int64
2   sin_month        12 non-null    float64
3   cos_month        12 non-null    float64
4   pred             12 non-null    float64
dtypes: float64(3), int64(2)
memory usage: 608.0 bytes
```

### 4 3. Modeling Temporal Data Mapped to Shark Attack Occurrences via Logistic Regression

#### Modeling Approach

#### Logistic Regression

When approaching how to model this problem, the nature of the what is being modeled was assessed. In this case, the outcome that was to be predicted was binary, whether a shark attack occurred or not based on the predictor at hand. In this case, logistic regression is a clear choice, as the predicted class in this method of supervised learning is a binary case.

## Accounting for Non-Shark Attack Occurrences (Negative Events)

After selecting logistic regression as the method of supervised learning to solve this problem, one main issue arose, which is the absence of negative cases. The data involved in this project is confirmed shark attack occurrences, and we do not have confirmed non-occurrences. This seems a silly distinction to make, as this is an obvious thought. There would be no point in recording non-shark attack occurrences, because they happen every day all around the world. The reason it is important is that the model will have to somehow account for these negative events.

To do so, the solution implemented in this particular project was the creation of synthetic negatives. Synthetic negatives allow one to account for all the non-shark attack occurrences. A reasonable question to ask would be: Wouldn't the synthetic negatives just represent the entire date-length of the dataset where attacks didn't occur (e.g. in this case, that would be every non-shark attack month, day, hour since 1931)? Unfortunately, one cannot feed a model this number of negative events, as shark attacks are such rare occurrences, the amount of non-attack occurrences are so large it greatly outweighs the attack occurrences. In this case, an equal length of data is sufficient to remove the model bias that would be had if one were to overwhelm it with negative events.

```
[39]: df_pos = df_clean[['date_at_location', 'time_of_attack_hour']].copy()
df_pos = df_pos[df_pos['time_of_attack_hour'] >= 0]
df_pos['attack_binary'] = 1

df_pos['month'] = df_pos['date_at_location'].dt.month

# df_pos['lunar_illumination_frac'] = df_pos['date_at_location'].
#   ↳ apply(moon_phase)
# df_pos['lunar_phase'] = df_pos['date_at_location'].apply(moon_phase_norm)

df_pos['lunar_illumination_phase'] = df_pos['date_at_location'].apply(moon_phase)
df_pos['lunar_illumination_frac'] = df_pos['lunar_illumination_phase'].
#   ↳ apply(illuminationfrac)
df_pos['lunar_phase'] = df_pos['lunar_illumination_phase'].apply(moon_phase_norm)

neg_date_range = pd.date_range(df_clean['date_at_location'].min(),
#   ↳ df_clean['date_at_location'].max(), freq = 'D')
neg_dates = np.random.choice(neg_date_range, len(df_pos), replace = True)

neg_hours = np.random.randint(0, 24, size = len(df_pos))

df_neg = pd.DataFrame({'date_at_location': neg_dates,
                        'time_of_attack_hour': neg_hours,
                        'attack_binary': 0 })

df_neg['month'] = df_neg['date_at_location'].dt.month
```



```

# df_neg['lunar_illumination_frac'] = df_neg['date_at_location'].
    ↳ apply(moon_phase)
# df_neg['lunar_phase'] = df_neg['lunar_illumination_frac'].
    ↳ apply(moon_phase_norm)

df_neg['lunar_illumination_phase'] = df_neg['date_at_location'].apply(moon_phase)
df_neg['lunar_illumination_frac'] = df_neg['lunar_illumination_phase'].
    ↳ apply(illuminationfrac)
df_neg['lunar_phase'] = df_neg['lunar_illumination_phase'].apply(moon_phase_norm)

df_model = pd.concat([df_pos, df_neg], ignore_index = True)

print(df_model.groupby('attack_binary').size())

```

```

attack_binary
0      574
1      574
dtype: int64

```

## 4.1 Correlation Heatmap for DF\_Model

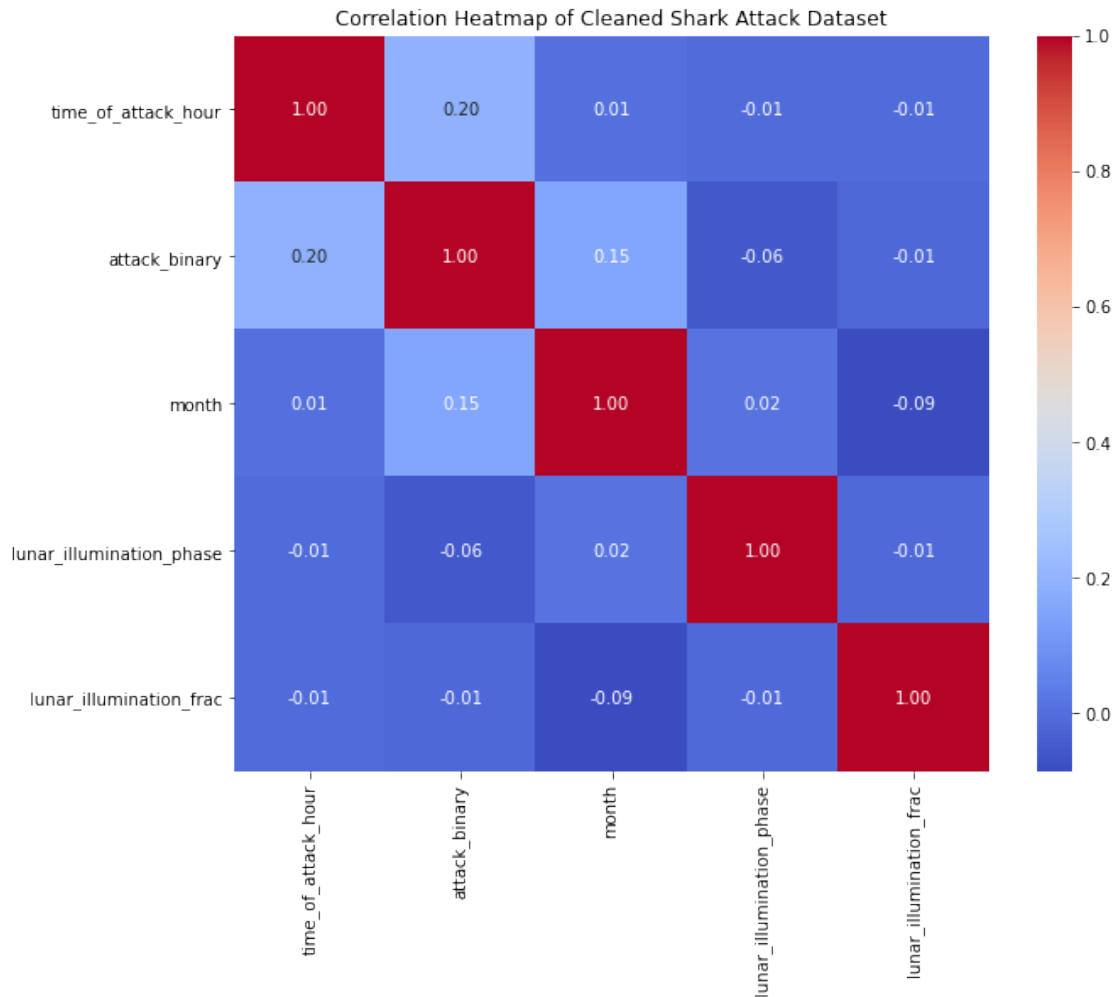
Now it looks like we have some significant results from our correlation heatmap with the synthetic negatives added in. There seems to be a correlation between attack binary (shark attack or not) and month, as well as hour. It should be noted that this is a point-biserial correlation (correlation between linear and binary values), so the correlation may not look as high as it actually is. The models below will prove out whether these are strong indicators or not

```

[40]: plt.figure(figsize=(10, 8))
sns.heatmap(df_model.select_dtypes(include = ["number"]).corr(),
            annot = True, cmap = "coolwarm", fmt = ".2f", cbar = True)

plt.title("Correlation Heatmap of Model Dataframe Shark Attack Data")
plt.show()

```



## 4.2 Logistic Regression: Shark Attack Occurrence ~ Lunar Illumination

As you can see, the results of our logistic regression below for lunar illumination on the attack occurrence were not strong. I would not say this is a strong predictor of attack occurrence

```
[74]: X = df_model[['lunar_illumination_frac']]
y = df_model['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.25,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
```

```

y_prob = model.predict_proba(X_test)[: , 1]

print(classification_report(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_prob))

print("Intercept is ", model.intercept_)
print("Coefficient is ", model.coef_)

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()

range_ = np.linspace(0,1,100).reshape(-1,1)
prob = model.predict_proba(range_)[: , 1]
plt.xlabel('Illumination Fraction')
plt.ylabel('Predicted Probability')
plt.title('Logreg Probability vs Illumination')
plt.ylim(0,1)
plt.plot(range_, prob, label = 'Attack Probability')
plt.legend()
plt.show()

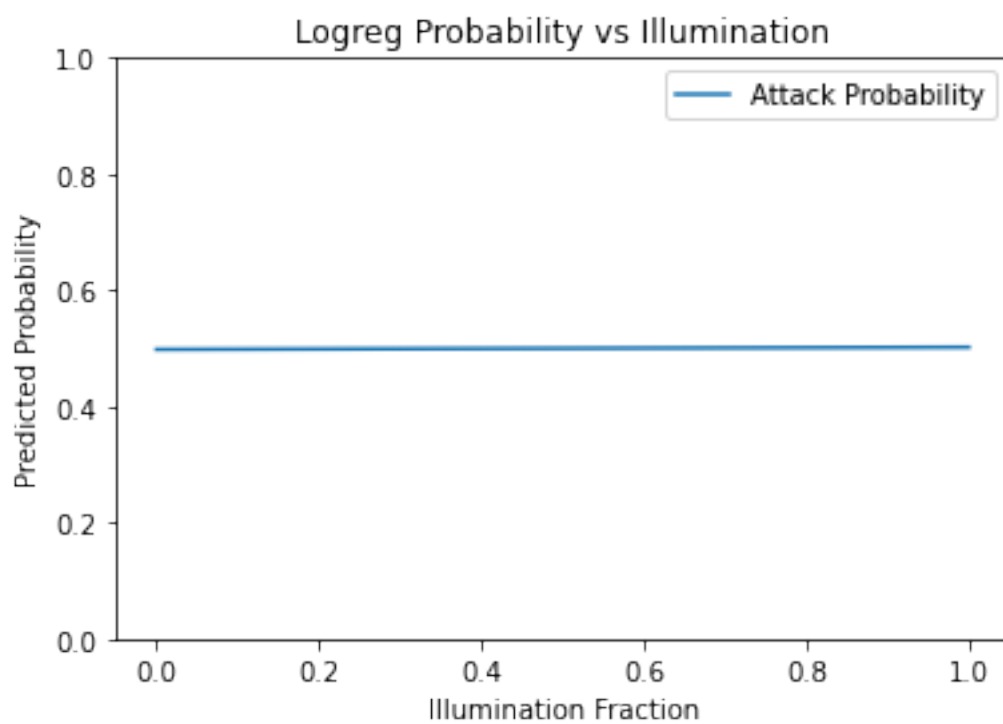
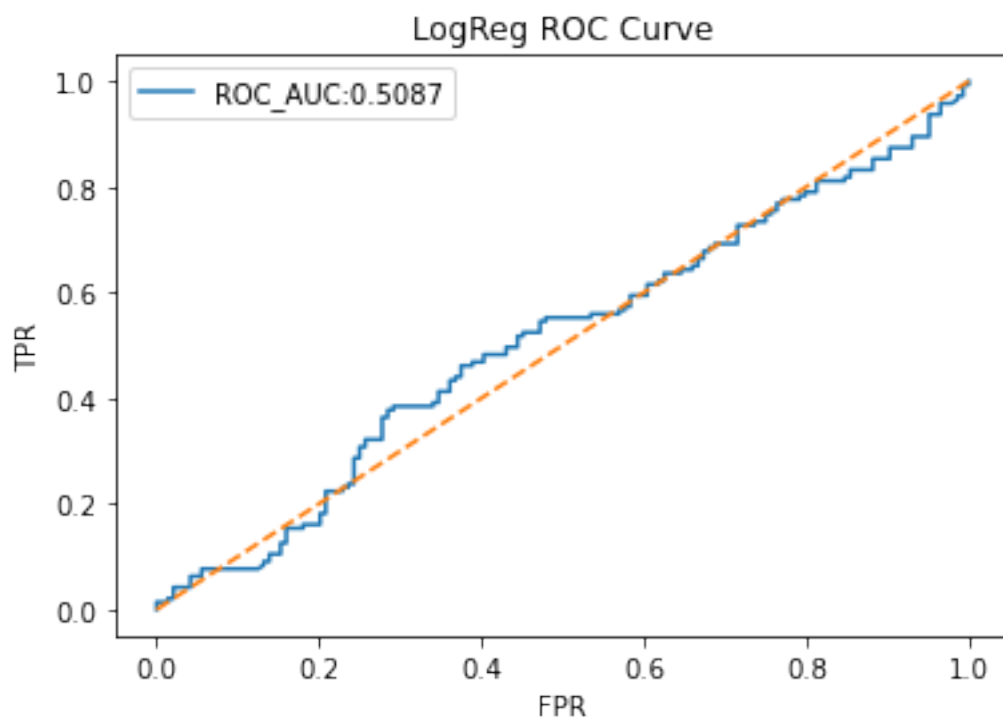
```

	precision	recall	f1-score	support
0	0.54	0.51	0.52	144
1	0.53	0.55	0.54	143
accuracy			0.53	287
macro avg	0.53	0.53	0.53	287
weighted avg	0.53	0.53	0.53	287

```

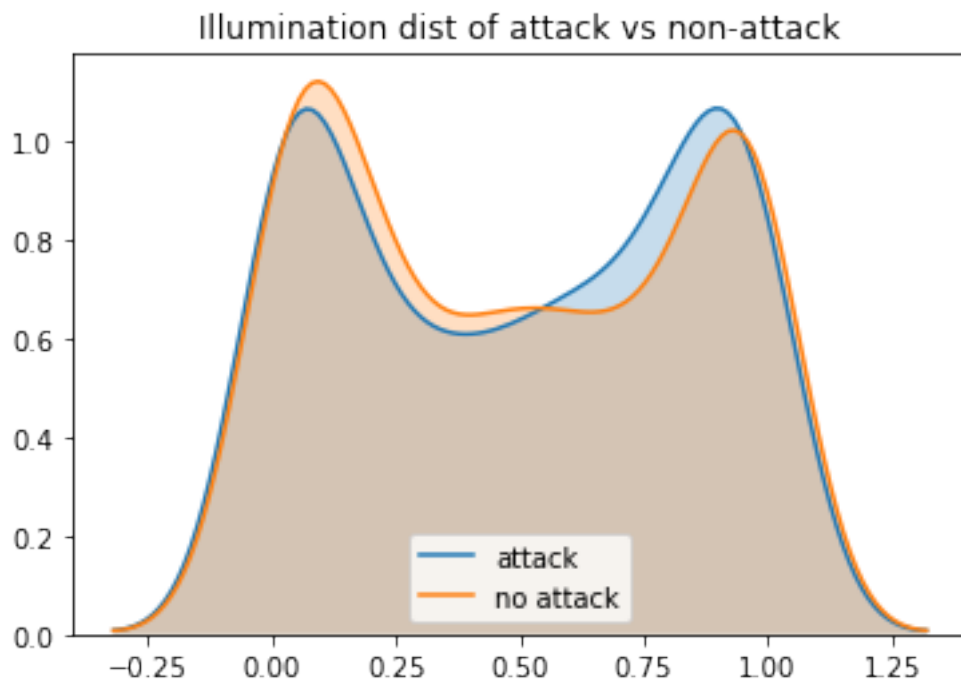
ROC_AUC: 0.5087412587412588
Intercept is [-0.00738919]
Coefficient is [[0.01505197]]

```



```
[75]: attack = df_model['attack_binary'] == 1
noattack = df_model['attack_binary'] == 0

sns.kdeplot(data = df_model.loc[attack, 'lunar_illumination_frac'], shade =_
↪True, label = 'attack')
sns.kdeplot(data = df_model.loc[noattack, 'lunar_illumination_frac'], shade =_
↪True, label = 'no attack')
plt.title("Illumination dist of attack vs non-attack")
plt.legend()
plt.show()
```



### 4.3 Logistic Regression: Shark Attack Occurrence ~ Month of Attack

Now looking at the monthly logistic regression model, this is a better indicator than lunar illumination. Looking at the predicted probabilities graph below, one can see that there is a higher predicted probability in the summer/fall months 6-10. You also can see that I had to use sin and cos encoding due to the cyclical nature of monthly data

```
[76]: df_model['month_sin'] = np.sin(2*np.pi*df_model['month']/12)
df_model['month_cos'] = np.cos(2*np.pi*df_model['month']/12)

X = df_model[['month_sin', 'month_cos']]
```

```

y = df_model['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.25,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1]

print(classification_report(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_prob))

print("Intercept is ", model.intercept_)
print("Coefficient is ", model.coef_)

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()

range_ = np.linspace(1,12,12).reshape(-1,1)
month_sin = np.sin(2*np.pi*range_/12)
month_cos = np.cos(2*np.pi*range_/12)
xp = np.c_[month_sin, month_cos]
prob = model.predict_proba(xp)[:,1]
plt.xlabel('Month')
plt.ylabel('Predicted Probability')
plt.title('Logreg Probability vs Month')
plt.ylim(0,1)
plt.plot(range_, prob, label = 'Attack Probability')
plt.legend()
plt.show()

```

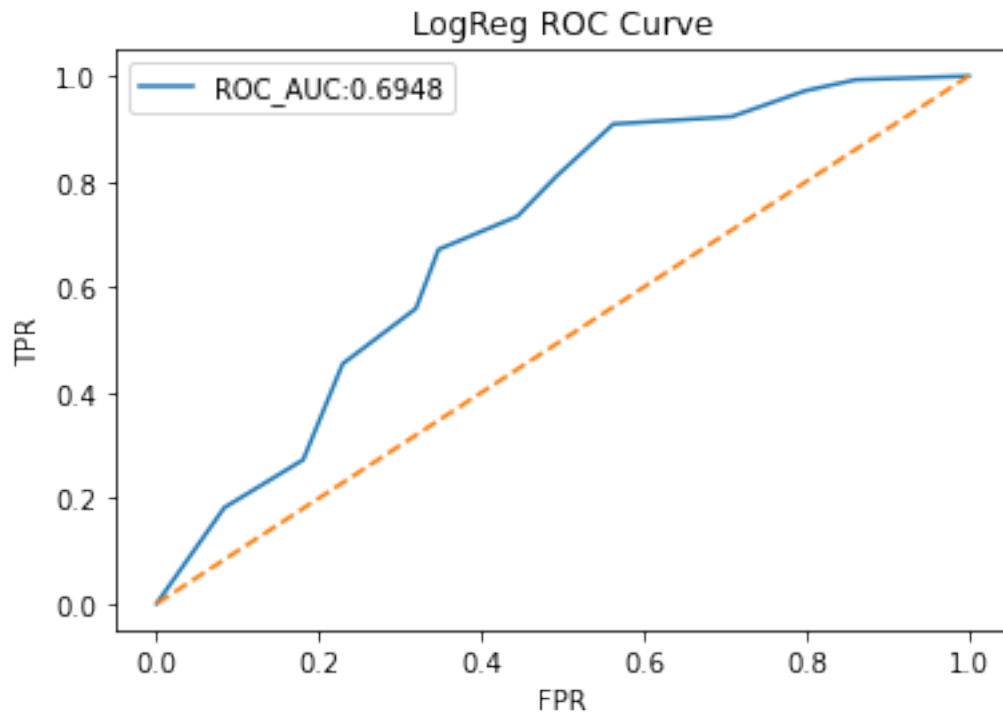
	precision	recall	f1-score	support
0	0.68	0.56	0.61	144
1	0.62	0.73	0.67	143
accuracy			0.64	287

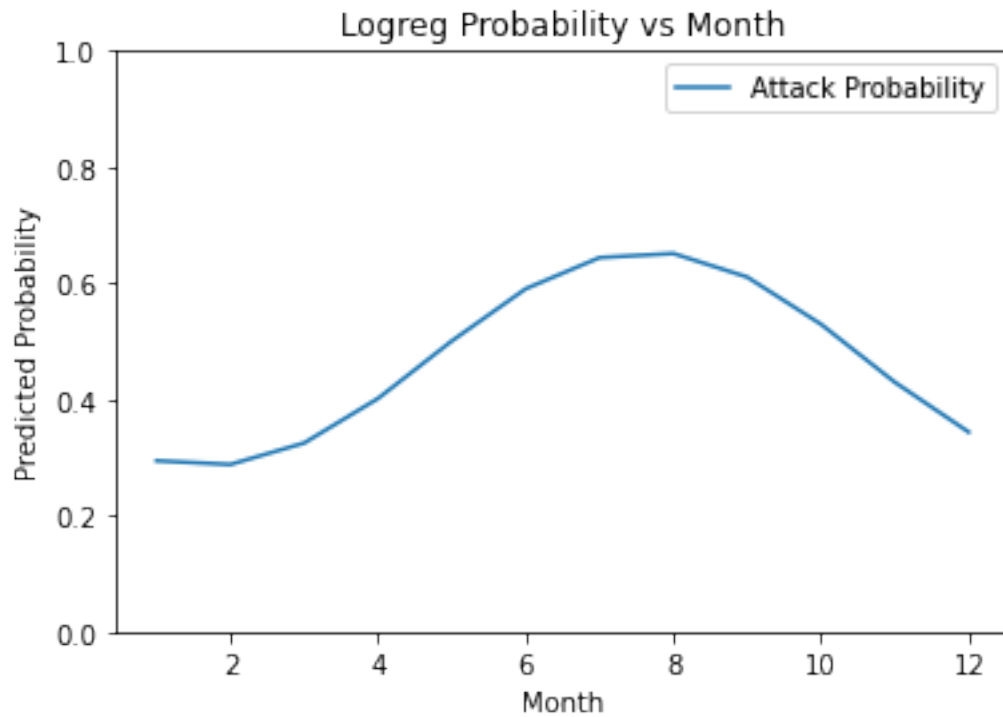
macro avg	0.65	0.64	0.64	287
weighted avg	0.65	0.64	0.64	287

ROC\_AUC: 0.6947843822843823

Intercept is [-0.14032613]

Coefficient is [[-0.59029376 -0.50607677]]

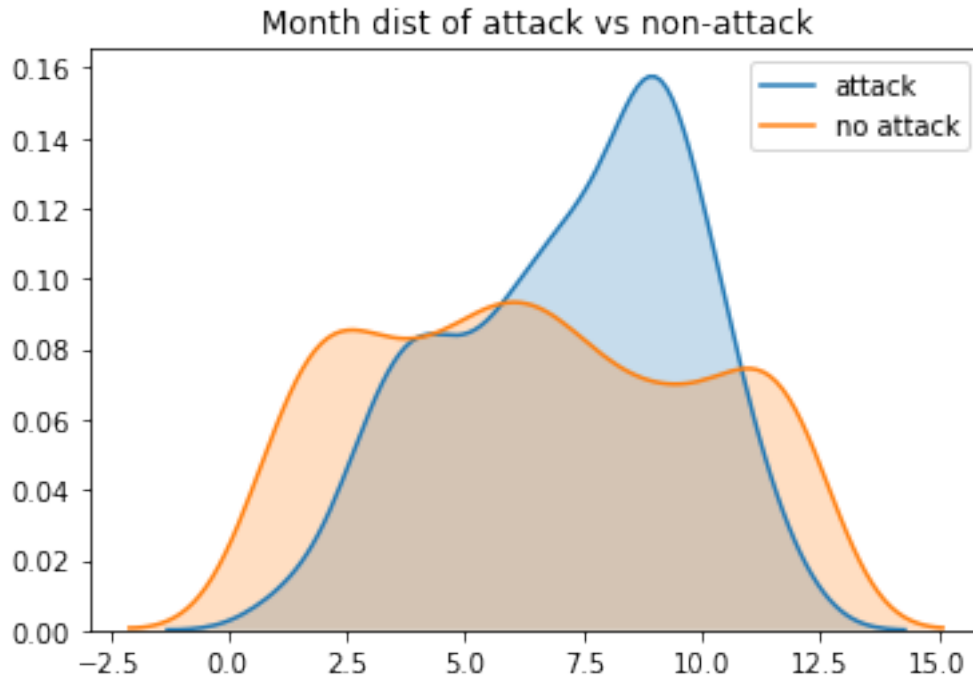




```
[77]: attack = df_model['attack_binary'] == 1
      noattack = df_model['attack_binary'] == 0

      sns.kdeplot(data = df_model.loc[attack, 'month'], shade = True, label = 'attack')
      sns.kdeplot(data = df_model.loc[noattack, 'month'], shade = True, label = 'noattack')
      plt.title("Month dist of attack vs non-attack")
      plt.legend()
      plt.show()
```





#### 4.4 Logistic Regression: Shark Attack Occurrence ~ Hour of Attack

As you can see, the hourly data did yield a high ROC AUC score and seems to be a strong indicator of shark attack occurrence. Unfortunately my probability predictions seem a bit inaccurate due to the fact that I used sin/cos encoding and we got an assumed two peaks, which is not accurate. This is because sin/cos enforces periodicity. Moving onto the next logistic regression, I decided to bucketize the hours to help with this accuracy.

```
[78]: df_model['hour_sin'] = np.sin(2*np.pi*df_model['time_of_attack_hour']/24)
df_model['hour_cos'] = np.cos(2*np.pi*df_model['time_of_attack_hour']/24)

X = df_model[['hour_sin', 'hour_cos']]
y = df_model['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.25,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[: , 1]

print(classification_report(y_test, y_pred))
```

```

print("ROC_AUC:", roc_auc_score(y_test, y_prob))

print("Intercept is ", model.intercept_)
print("Coefficient is ", model.coef_)

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()

range_ = np.arange(0,24)
hour_sin = np.sin(2*np.pi*range_/12)
hour_cos = np.cos(2*np.pi*range_/12)
xp = np.c_[hour_sin, hour_cos]
prob = model.predict_proba(xp)[: ,1]
plt.xlabel('Hour')
plt.ylabel('Predicted Probability')
plt.title('Logreg Probability vs Hour')
plt.ylim(0,1)
plt.plot(range_, prob, label = 'Attack Probability')
plt.legend()
plt.show()

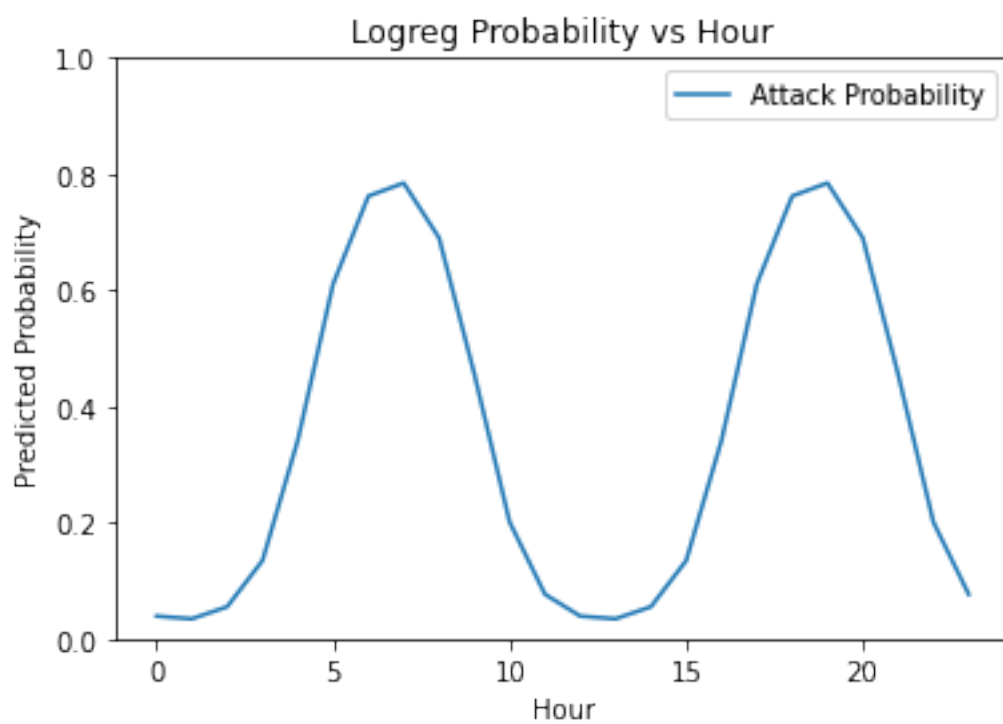
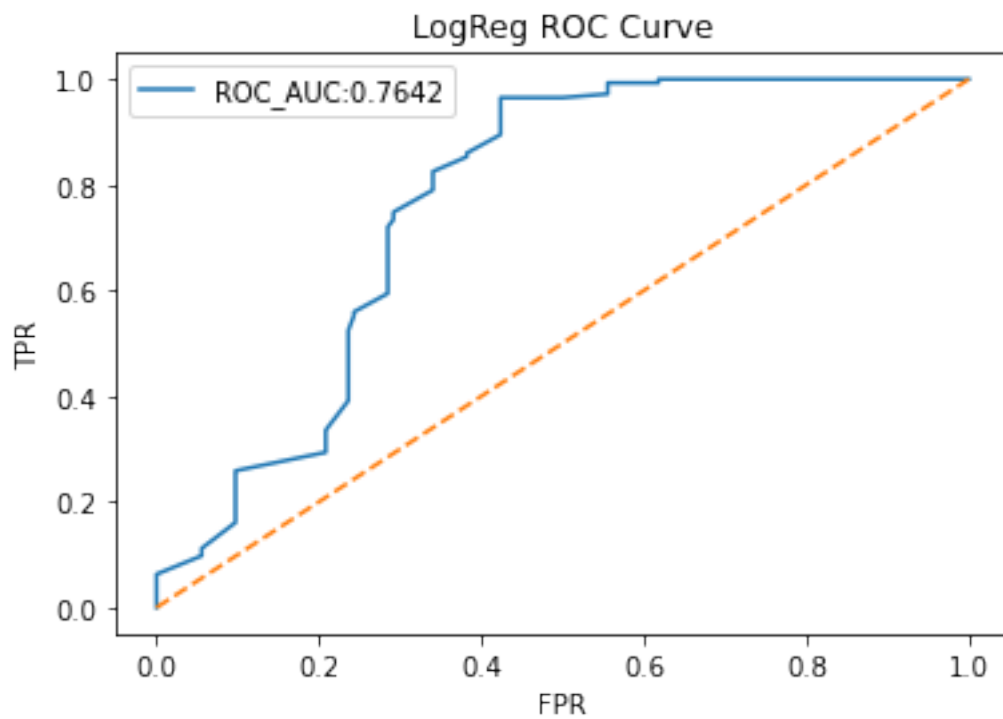
```

	precision	recall	f1-score	support
0	0.79	0.66	0.72	144
1	0.70	0.82	0.76	143
accuracy			0.74	287
macro avg	0.74	0.74	0.74	287
weighted avg	0.75	0.74	0.74	287

```

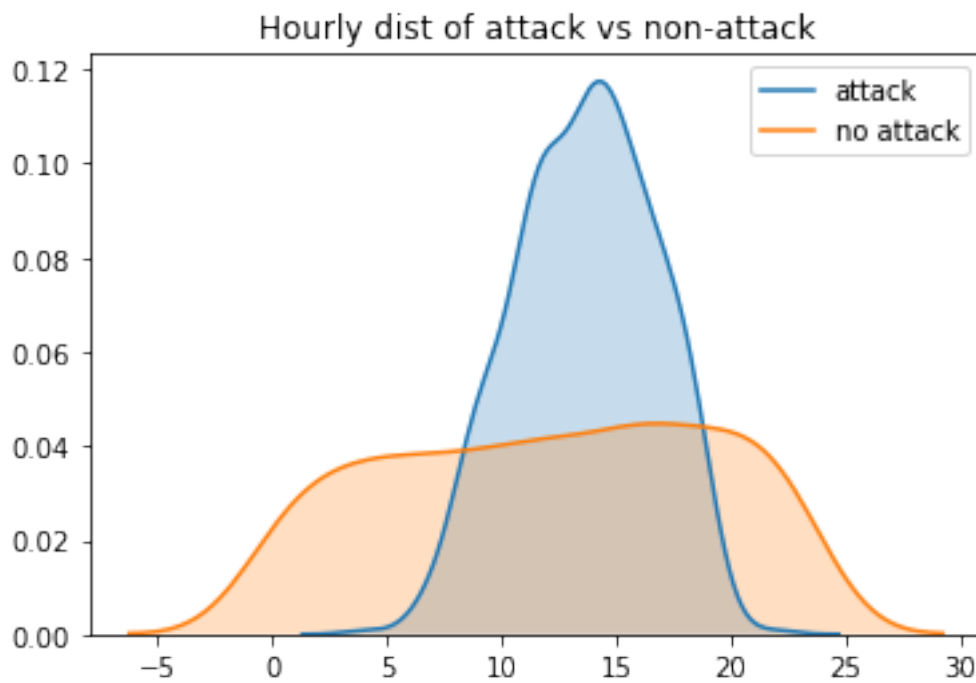
ROC_AUC: 0.7642288267288266
Intercept is [-1.01823709]
Coefficient is [[-0.84152096 -2.18175898]]

```



```
[79]: attack = df_model['attack_binary'] == 1
noattack = df_model['attack_binary'] == 0

sns.kdeplot(data = df_model.loc[attack, 'time_of_attack_hour'], shade = True,
            ↪label = 'attack')
sns.kdeplot(data = df_model.loc[noattack, 'time_of_attack_hour'], shade = True,
            ↪label = 'no attack')
plt.title("Hourly dist of attack vs non-attack")
plt.legend()
plt.show()
```



#### 4.5 Logistic Regression: Shark Attack Occurrence ~ Hourly Bucket (Time of Day)

For this model, I bucketed the hours into 5 different categories: Night, morning, mid-day, afternoon, evening, night. As you can see, I split the day categories into 4 groups evenly, and then had the night as a large bucket by itself that is 2x the size as the other buckets. The reasoning behind this is because there are very few beachgoers in the evening, so it doesn't make sense for it to have the same weight as the other categories.

```
[80]: def new_hour_bucket(hr):
        if 5 <= hr < 9:
```

```

        return 'morning'
    if 9<= hr <13:
        return 'mid-day'
    if 13<= hr <17:
        return 'afternoon'
    if 17<= hr <21:
        return 'evening'
    else:
        return 'night'

df_model['hour_bucket'] = df_model['time_of_attack_hour'].apply(new_hour_bucket)

X = pd.get_dummies(df_model[['hour_bucket']], drop_first=True)
y = df_model['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.25,
    ↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1]

print(classification_report(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_prob))

print("Intercept is ", model.intercept_)
print("Coefficient is ", model.coef_)

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()

prob = model.predict_proba(X)[: ,1]
df_plot = pd.DataFrame({'bucket': df_model['hour_bucket'], 'prob':prob})
df_plot.groupby('bucket').mean().reindex(['night', 'morning', 'mid-day',
    ↳'afternoon', 'evening']).plot(kind = 'bar')

```

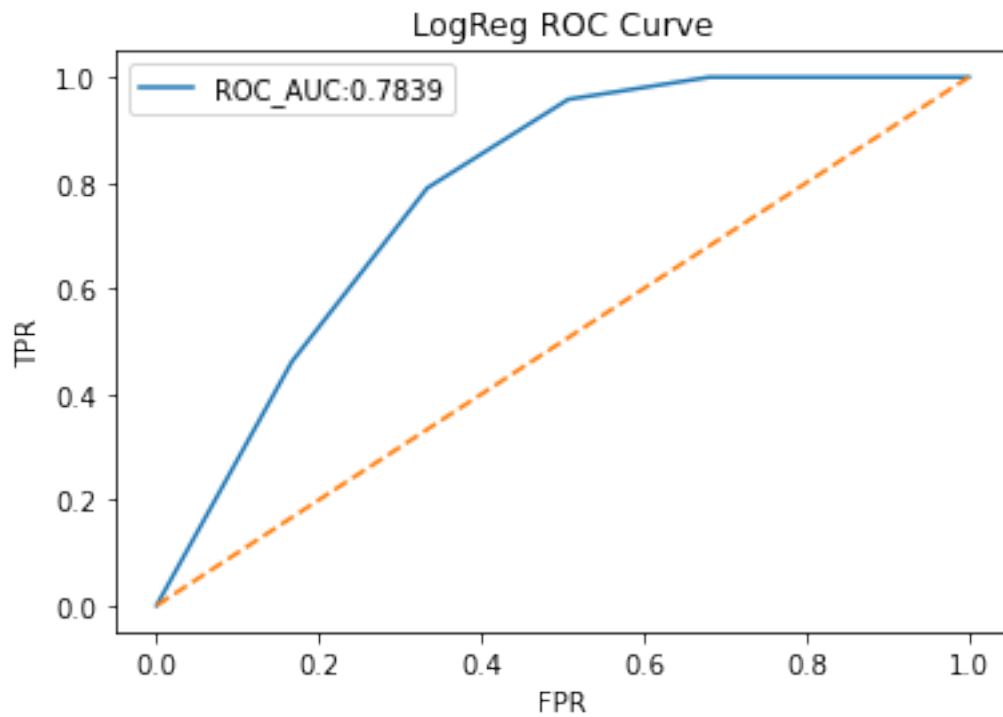
```
plt.xlabel('Hour Bucket')
plt.ylabel('Predicted Probability')
plt.show()
```

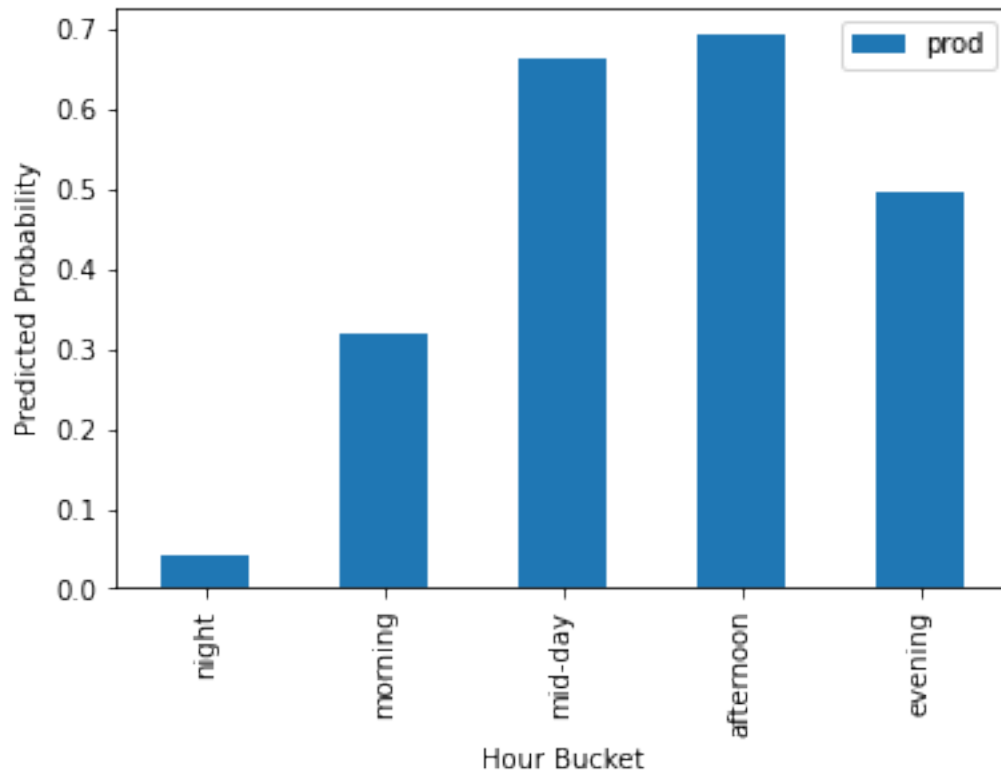
	precision	recall	f1-score	support
0	0.76	0.67	0.71	144
1	0.70	0.79	0.74	143
accuracy			0.73	287
macro avg	0.73	0.73	0.73	287
weighted avg	0.73	0.73	0.73	287

ROC\_AUC: 0.7839452214452214

Intercept is [0.81176111]

Coefficient is [[-0.83249805 -0.13858556 -1.56351051 -3.98673338]]





#### 4.6 Logistic Regression: Shark Attack Occurrence ~ Lunar Phase

Since I had some success bucketting the hourly data, I am going to try and do the same with the lunar illumination data, and used my already bucketted version of this: the Lunar Phases (First Quarter, Full, Last Quarter, New).

As you can see below, the results were quite insignificant. Although we saw a slight skew toward the New Moon in our EDA, you can see below that the Logistic Regression did not confirm this difference to be significant.

```
[81]: X = pd.get_dummies(df_model[['lunar_phase']], drop_first=True)
y = df_model['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[: , 1]
```

```

print(classification_report(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_prob))

print("Intercept is ", model.intercept_)
print("Coefficient is ", model.coef_)

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()

prob = model.predict_proba(X)[: ,1]
df_plot = pd.DataFrame({'Lunar Phase': df_model['lunar_phase'], 'prob':prob})
df_plot.groupby('Lunar Phase').mean().reindex(['first_quarter', 'full', 'last_quarter', 'new']).plot(kind = 'bar')
plt.xlabel('Lunar Phase')
plt.ylabel('Predicted Probability')
plt.show()

```

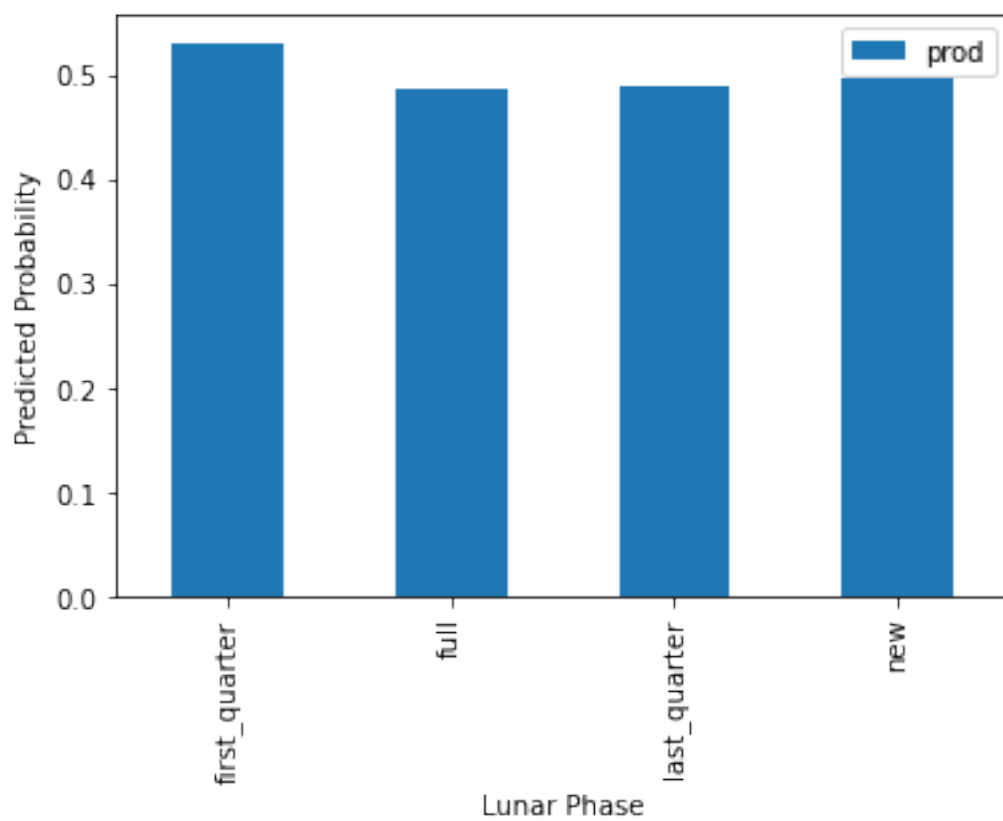
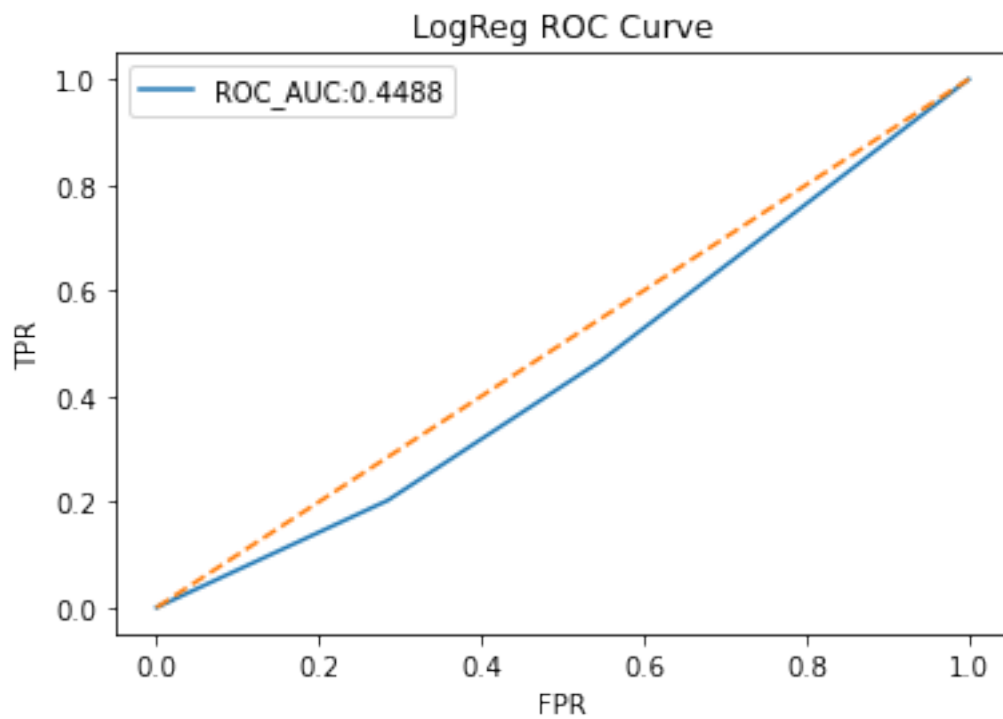
	precision	recall	f1-score	support
0	0.47	0.72	0.57	144
1	0.41	0.20	0.27	143
accuracy			0.46	287
macro avg	0.44	0.46	0.42	287
weighted avg	0.44	0.46	0.42	287

```

ROC_AUC: 0.4487665112665113
Intercept is [0.12147801]
Coefficient is [[-0.1751672 -0.16682189 -0.13025515]]

```





## 4.7 Logistic Regression: Shark Attack Occurrence ~ Month of Attack + Hourly Bucket (Time of Attack)

For this final logistic regression model, I wanted to combine Month and the hourly bucket data to see if we could get a significant result. Because those two features alone, showed some promising results, I decided to combine them to see if I could create an enhanced model. The ROC\_AUC score was quite significant for this value, so I think this model did yield a significant result, showing both together are strong predictors of an attack occurrence. More discussion on this significant result will be had in the Conclusion section

```
[82]: X = pd.get_dummies(df_model[['month', 'hour_bucket']], drop_first=True)
y = df_model['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1]

print(classification_report(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_prob))

print("Intercept is ", model.intercept_)
print("Coefficient is ", model.coef_)

fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()

df_model['prob'] =model.predict_proba(X)[: ,1]

bucket_ = df_model.groupby('hour_bucket')['prob'].mean()
plt.figure()
```

```

bucket_.reindex(['night', 'morning', 'mid-day', 'afternoon', 'evening']).
    ↪plot(kind = 'bar')
plt.ylim(0,1)
plt.xlabel('Hour Bucket')
plt.ylabel('Predicted Probability')
plt.show()

month_ = df_model.groupby('month')['prob'].mean()
plt.figure
month_.plot(marker = 'o')
plt.ylim(0,1)
plt.xlabel('Month')
plt.ylabel('Predicted Probability')
plt.show()

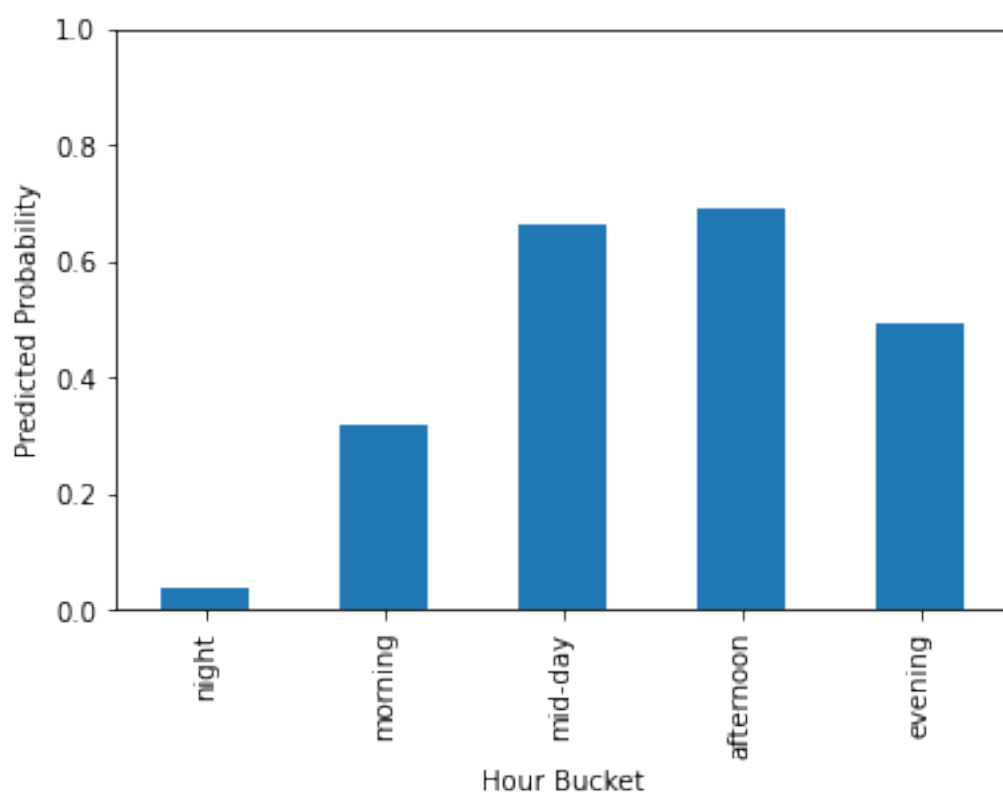
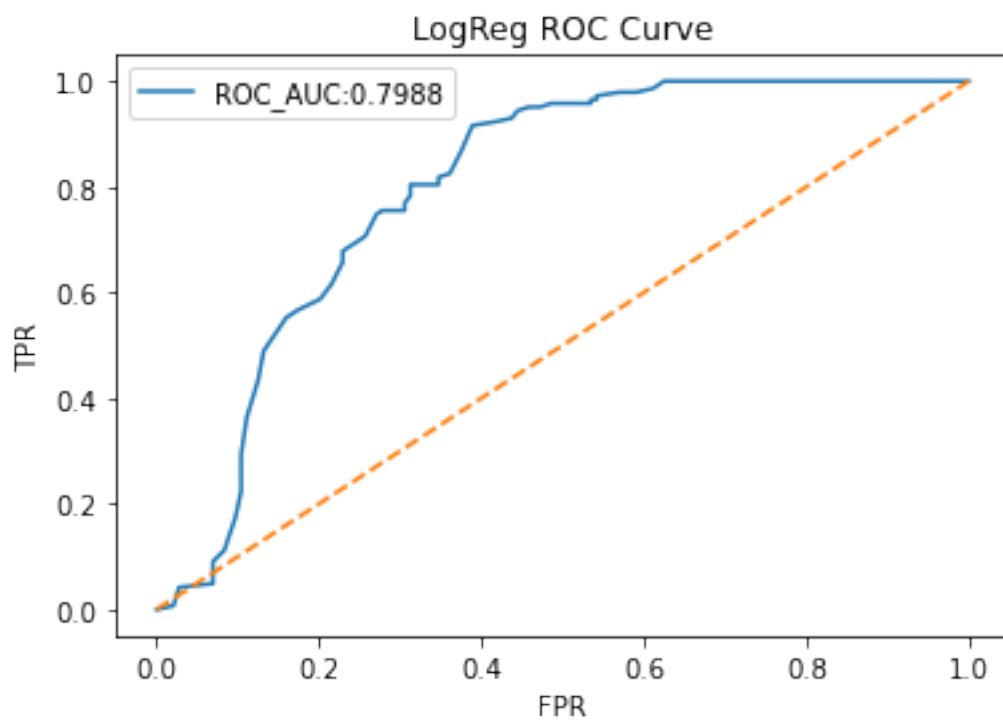
```

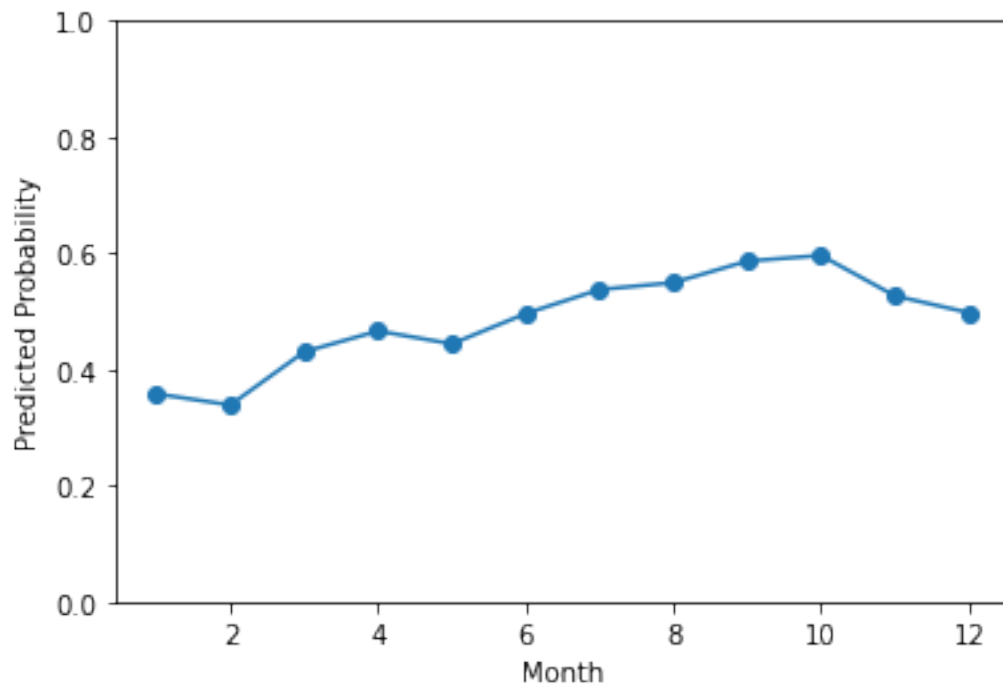
	precision	recall	f1-score	support
0	0.88	0.61	0.72	144
1	0.70	0.92	0.79	143
accuracy			0.76	287
macro avg	0.79	0.76	0.76	287
weighted avg	0.79	0.76	0.76	287

ROC\_AUC: 0.7988053613053613

Intercept is [0.25151363]

Coefficient is [[ 0.08329165 -0.86017853 -0.14383528 -1.60011243 -4.00051617]]





## 5 4. Modeling Pressure and Temperature on Attack Occurrence

I got some interesting results with the Temporal data, but now I would like to take the analysis a step further and bring in some weather features, such as pressure and temperature. Below I am pulling in data from meteostat (a weather data library in Python with NOAA station weather data) and merging it in with my original dataframe df\_clean

### 5.1 Meteostat Data Import and Cleansing

```
[36]: # !pip uninstall -y meteostat
!pip install -q "meteostat==1.6.1"
!pip install "pandas>=1.2,<1.4"
```

```
[notice] A new release of pip is
available: 23.3.2 -> 24.0
[notice] To update, run:
pip install --upgrade pip
Requirement already satisfied: pandas<1.4,>=1.2 in
/opt/conda/lib/python3.7/site-packages (1.3.5)
Requirement already satisfied: python-dateutil>=2.7.3 in
```

```
/opt/conda/lib/python3.7/site-packages (from pandas<1.4,>=1.2) (2.8.1)
Requirement already satisfied: pytz>=2017.3 in /opt/conda/lib/python3.7/site-
packages (from pandas<1.4,>=1.2) (2020.1)
Requirement already satisfied: numpy>=1.17.3 in /opt/conda/lib/python3.7/site-
packages (from pandas<1.4,>=1.2) (1.18.4)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-
packages (from python-dateutil>=2.7.3->pandas<1.4,>=1.2) (1.14.0)
```

```
[notice] A new release of pip is
available: 23.3.2 -> 24.0
```

```
[notice] To update, run:
pip install --upgrade pip
```

```
[37]: from meteostat import Point, Hourly, Stations
```

After importing all data and libraries needed, I am printing all the stations below to see which ones I want to include in my dataset. Ideally I would like bring in all counties that are in `df_clean`, but because of time and length considerations, I will have to save that for future analysis.

Out of the three counties I see below, I believe my best options are Volusia, Brevard and Duval. All three represent three distinct Atlantic Coast regions in Florida (South, Central, and North Coast), and have a large number of datapoints for me to analyze.

```
[38]: print(df_clean['county'].value_counts(dropna = False))
```

```
volusia county      322
brevard county      104
palm beach county   66
st. johns county    31
st. lucie county     31
duval county         29
martin county        27
indian river county  19
miami-dade county   14
broward county       13
flagler county        6
nassau county         6
dade county          2
miami-dade           1
Name: county, dtype: int64
```

Here I am printing all the stations that are nearby the county seat of all these counties.

```
[39]: volusia = Stations().nearby(29.0283, -81.0755)
brevard = Stations().nearby(28.2639, -80.7214)
duval = Stations().nearby(30.3500, -81.6035)

volusia = volusia.inventory('hourly')
brevard = brevard.inventory('hourly')
```

```

duval = duval.inventory('hourly')

vstations = volusia.fetch(10)
bstations = brevard.fetch(10)
dstations = duval.fetch(10)

print(vstations[['name', 'latitude', 'longitude', 'region']])
print(bstations[['name', 'latitude', 'longitude', 'region']])
print(dstations[['name', 'latitude', 'longitude', 'region']])

```

	name	latitude	longitude	\
id				
KEVB0	New Smyrna Beach / Isleboro	29.0557	-80.9489	
74787	Daytona Beach / Mansfield Mobile Home Park	29.1799	-81.0581	
KDED0	Deland / Orangewood Mobile Home Park	29.0670	-81.2837	
KOMNO	Ormond Beach / Bear Creek Mobile Home Park	29.3006	-81.1136	
KSFB0	Sanford / Orlando / Midway	28.7767	-81.2355	
KXFLO	Palm Coast / Palm Terrace Mobile Manor	29.4674	-81.2063	
KTTS0	Cape Kennedy / Wilson	28.6149	-80.6944	
KORLO	Orlando / Azalea Park	28.5455	-81.3329	
KTIX0	Titusville / Manatee Hammock Park Mobile Home ...	28.5148	-80.7992	
72205	Orlando Airport	28.4167	-81.0000	

	region
id	
KEVB0	FL
74787	FL
KDED0	FL
KOMNO	FL
KSFB0	FL
KXFLO	FL
KTTS0	FL
KORLO	FL
KTIX0	FL
72205	FL

	name	latitude	longitude	\
id				
74795	Cocoa / Patrick Air Force Base	28.2333	-80.6000	
72204	Melbourne International	28.1000	-80.6500	
74794	Cocoa Beach, Cape Canaveral Air Force Station ...	28.4667	-80.5667	
KTIX0	Titusville / Manatee Hammock Park Mobile Home ...	28.5148	-80.7992	
72205	Orlando Airport	28.4167	-81.0000	
KTTS0	Cape Kennedy / Wilson	28.6149	-80.6944	
YZNBG	Sebastian Municipal Airport	27.8132	-80.4956	
KORLO	Orlando / Azalea Park	28.5455	-81.3329	
KISMO	Orlando / Country Life Family Mobile Home Park	28.2898	-81.4371	
74793	Gifford	27.6561	-80.4181	

```

region
id
74795      FL
72204      FL
74794      FL
KTIx0      FL
72205      FL
KTTS0      FL
YZNBG      FL
KORLO      FL
KISMO      FL
74793      FL

```

```

name latitude longitude \
id
74782      Jacksonville / Hatch Road Mobile Home Park 30.3363 -81.5144
NFUKD      Jacksonville Naval Air Station 30.2347 -81.6746
KNRB0      Mayport 30.3914 -81.4245
KHEG0      Jacksonville / Normandy Estates Mobile Home Co... 30.2778 -81.8059
KNEN0      Jackson / Mandeville 30.3502 -81.8832
KVQQ0      Jacksonville / Baileys Mobile Home Park 30.2187 -81.8767
KFHB0      Fernandina Beach / Amelia Island 30.6118 -81.4612
KSGJ0      St Augustine / Kingsley Mobile Home Park 29.9593 -81.3397
K2CB0      Camp Blanding Mil Res(Starke) / Kingsley Beach 29.9669 -81.9832
72206      Jacksonville International 30.5000 -81.0333

```

```

region
id
74782      FL
NFUKD      FL
KNRB0      FL
KHEG0      FL
KNEN0      FL
KVQQ0      FL
KFHB0      FL
KSGJ0      FL
K2CB0      FL
72206      FL

```

I have chosen the following stations based on their proximity to beaches (where shark attacks happen)

Volusia: KEVB0 New Smyrna Beach / Isleboro 29.0557 -80.9489  
 Brevard: 74795 Cocoa / Patrick Air Force Base 28.2333 -80.6000  
 Duval: KNRB0 Mayport 30.3914 -81.4245

```
[40]: for st_id in ['KEVB0', '74795', 'KNRB0']:
```



```

data = Hourly(st_id, start = pd.Timestamp("2003-04-08"), end = pd.
↳Timestamp("2003-04-08")).fetch()
print(st_id)
print(data.columns)
print(data[['temp', 'pres']].head())

```

```

KEVB0
Index(['temp', 'dwpt', 'rhum', 'prcp', 'snow', 'wdir', 'wspd', 'wpgt', 'pres',
      'tsun', 'coco'],
      dtype='object')
Empty DataFrame
Columns: [temp, pres]
Index: []
74795
Index(['date', 'hour', 'temp', 'dwpt', 'rhum', 'prcp', 'snow', 'wdir', 'wspd',
      'wpgt', 'pres', 'tsun', 'coco'],
      dtype='object')
Empty DataFrame
Columns: [temp, pres]
Index: []
KNRB0
Index(['temp', 'dwpt', 'rhum', 'prcp', 'snow', 'wdir', 'wspd', 'wpgt', 'pres',
      'tsun', 'coco'],
      dtype='object')
      temp    pres
time
2003-04-08  22.8  1017.5

```

Now I am going through and creating a new dataframe that will hold all the hourly meteorological data I want, and then merging it with my shark attack dataset (df\_clean).

I should also note here, that I am only taking data from 2001 onward, because the weather data prior to this time can be a bit spotty.

Additionally I am creating 8 new features from the data:

t\_delta12- The temperature change from 12 hours past to the time of occurrence t\_delta24- The temperature change from 24 hours past to the time of occurrence p\_delta12- The pressure change from 12 hours past to the time of occurrence p\_delta24- The pressure change from 24 hours past to the time of occurrence

t\_slope12- The slope of the temperature change from 12 hours past to the time of occurrence t\_slope24- The slope of the temperature change from 24 hours past to the time of occurrence p\_slope12- The slope of the pressure change from 12 hours past to the time of occurrence p\_slope24- The slope of the pressure change from 24 hours past to the time of occurrence

I added these features because I was interested to see if not only the pressure and temperature in itself were strong indicators of an attack occurrence, but I also wanted to understand if the change in these features could be of significance too.

```

[41]: county_stat = {'volusia county': 'KEVB0', 'brevard county': '74795', 'duval_
      ↪county': 'KNRB0'}

df_a = df_clean[df_clean['county'].str.contains("Volusia|Brevard|Duval", case =_
      ↪False)].copy()

df_meteo = []

for county, st_id in county_stat.items():
    subset = df_a[df_a['county'].str.contains(county, case = False)].copy()
    subset = subset[subset['date_at_location'] >= pd.Timestamp("2000-01-01") ]

    if subset.empty:
        continue

    start = subset['date_at_location'].min()
    end = subset['date_at_location'].max()

    meteo = Hourly(st_id, start, end).fetch()
    meteo = meteo.sort_index()

    meteo['t_delta12'] = meteo['temp'].diff(12)
    meteo['t_delta24'] = meteo['temp'].diff(24)
    meteo['p_delta12'] = meteo['pres'].diff(12)
    meteo['p_delta24'] = meteo['pres'].diff(24)

    meteo['t_slope12'] = (meteo['temp'] - meteo['temp'].shift(12))/12
    meteo['t_slope24'] = (meteo['temp'] - meteo['temp'].shift(24))/24
    meteo['p_slope12'] = (meteo['pres'] - meteo['pres'].shift(12))/12
    meteo['p_slope24'] = (meteo['pres'] - meteo['pres'].shift(24))/24

    subset['attack_hour'] = subset['date_at_location'].dt.floor('H')
    meteo_subset = pd.merge(subset, meteo, left_on = 'attack_hour', right_index_
    ↪= True, how = 'left')
    meteo_subset['county_match'] = county
    df_meteo.append(meteo_subset)

df_meteo = pd.concat(df_meteo, ignore_index = True)

print(df_meteo[['county_match', 'date_at_location', 'temp', 'pres', 't_delta12',
      't_delta24', 'p_delta12', 'p_delta24']].head(200))

```

	county_match	date_at_location	temp	pres	t_delta12	t_delta24	\
0	volusia county	2000-06-02	NaN	NaN	NaN	NaN	
1	volusia county	2000-06-19	NaN	NaN	NaN	NaN	
2	volusia county	2000-07-02	NaN	NaN	NaN	NaN	
3	volusia county	2000-07-02	NaN	NaN	NaN	NaN	

4	volusia county	2000-07-04	NaN	NaN	NaN	NaN
..	...	...	...	...	...	...
195	volusia county	2021-08-14	26.0	1020.0	-3.0	-2.0
196	volusia county	2021-09-05	26.0	1014.0	0.0	0.0
197	volusia county	2021-09-08	27.0	1015.0	1.0	0.0
198	volusia county	2021-09-09	24.0	1013.0	-3.0	-3.0
199	volusia county	2021-09-11	26.0	1016.0	1.0	0.0

	p_delta12	p_delta24
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
..	...	...
195	1.0	1.0
196	-1.0	-1.0
197	1.0	1.0
198	-2.0	-2.0
199	4.0	4.0

[200 rows x 8 columns]

Below I am inspecting my final dataframe to see what I have. It looks as though I have a lot of nan values for air temp and pressure. To combat this, I will use bfill and ffill, which will help me backfill (to the nearest hour) any missing weather data.

```
[42]: print(len(df_meteo))
print(df_meteo.isna().sum())
#print((df_meteo.isna().mean()*100).round(2))
```

284	
attack_classification	0
outcome	0
county	0
date_at_location	0
Year	0
time_of_attack	2
victim_activity	0
sunlight_conditions	0
gen_weather	0
air_temp	134
water_temp	99
closest_phase	0
waxing_waning	0
common_name	0
provocative_acts	0
activity_addenda	0

outcome_binary	0
outcome_ternary	0
time_of_attack_hour	34
time_unknown_flag	0
time_bucket	0
month	0
season	0
hour_sin	34
hour_cos	34
lunar_illumination_phase	0
lunar_illumination_frac	0
lunar_illumination_pct	0
lunar_phase	0
attack_hour	0
temp	188
dwpt	188
rhum	188
prcp	234
snow	284
wdir	191
wspd	189
wpgt	284
pres	207
tsun	284
coco	264
t_delta12	189
t_delta24	189
p_delta12	211
p_delta24	211
t_slope12	189
t_slope24	189
p_slope12	209
p_slope24	211
county_match	0
dtype:	int64

```
[43]: df_meteo['p_nan'] = df_meteo['pres'].isna().astype(int)
df_meteo['t_nan'] = df_meteo['temp'].isna().astype(int)

df_meteo2 = df_meteo.groupby('county_match').apply(lambda g: g.ffill().bfill()).
↳reset_index(drop = True)
#print(df_meteo2.isna().sum())

df_meteo2['p_nan'] = df_meteo['p_nan'].values
df_meteo2['t_nan'] = df_meteo['t_nan'].values
```

```
[44]: print(df_meteo2[['temp', 'pres', 't_delta12', 'p_delta12']].describe())
print(df_meteo2['county_match'].value_counts())
```

	temp	pres	t_delta12	p_delta12
count	284.000000	284.000000	284.000000	284.000000
mean	25.816549	1015.94507	2.089437	-1.398239
std	2.823145	1.91828	4.898622	1.735280
min	12.000000	1005.00000	-11.000000	-8.200000
25%	24.000000	1016.00000	-1.000000	-2.000000
50%	26.000000	1016.00000	2.000000	-2.000000
75%	29.000000	1016.00000	6.000000	-1.100000
max	29.400000	1022.00000	13.000000	13.200000
volusia county	221			
brevard county	45			
duval county	18			

Name: county\_match, dtype: int64

## 5.2 EDA on Pressure and Temperature

### Weather Visualizations

The below visualizations are some significant features that were found to be of note when inspecting the weather data. Data was pulled in from meteostat, which is a python library that contains weather data pulled from NOAA weather stations. With this, a new data byproduct was created, enriching the original shark attack data with meteorological features. A further note regarding this byproduct is that non-attack occurrences were created for comparison to the attack occurrences, but this will be discussed later in this paper.

There are some basic box-and-whisker plots that show the distribution of attacks across three different features: temperature, change in pressure, and rate of change (or slope) of pressure. Other features were visualized, but these were the most pertinent to modeling. As one can see by the below visualizations, there is a slight skew toward higher temperature on the attack values vs non-attack values. Additionally, the change in pressure and slope of pressure change give an even slighter skew higher. Moving to the modeling phase, these observations will be tested for their statistical significance.

As you can see, just looking at the pressure and time data alone, to identify shark attack occurrence, doesnt do much, so I am going to do something similar to what I did before for my Temporal models and create a synthetic negative set that will represent a non-attack occurrence.

```
[45]: # sns.distplot(df_meteo2['temp'], bins = 20, kde = True)
# plt.show()

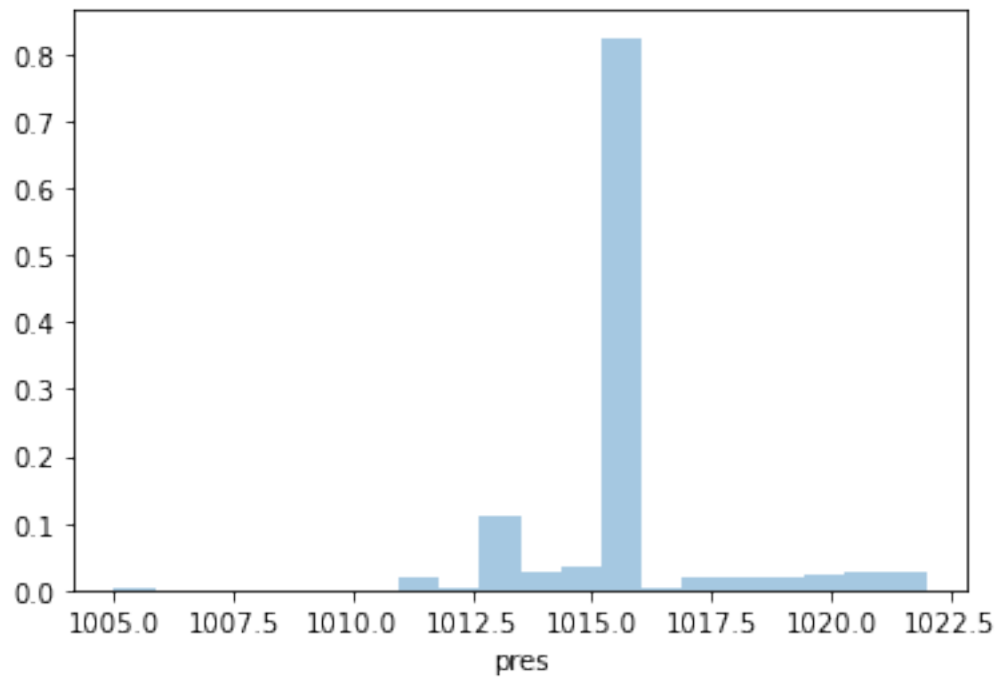
sns.distplot(df_meteo2['pres'], bins = 20, kde = True)
plt.show()

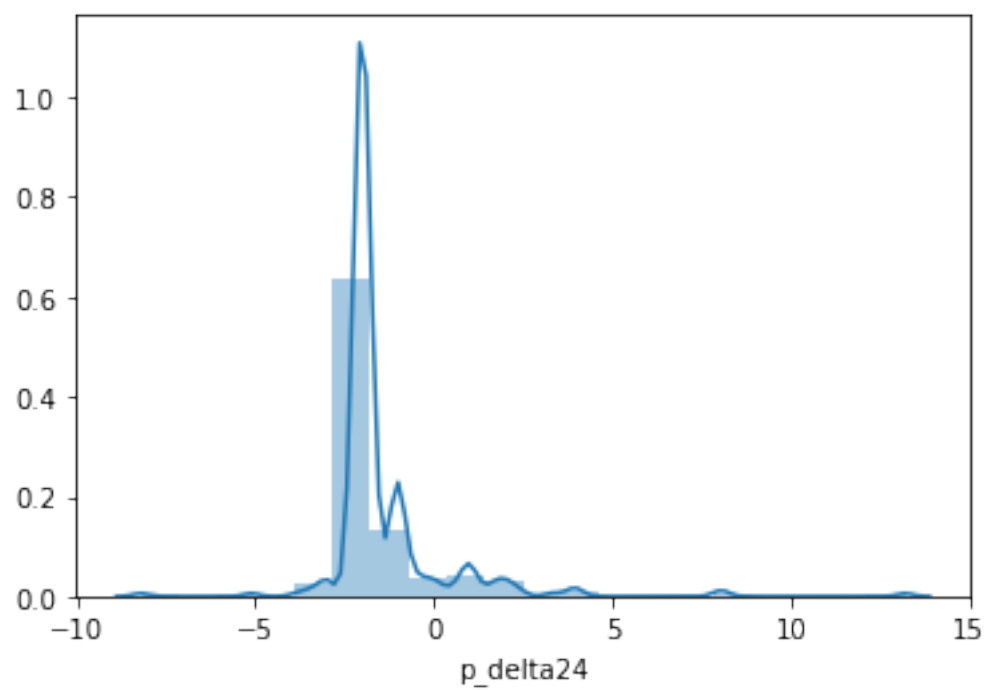
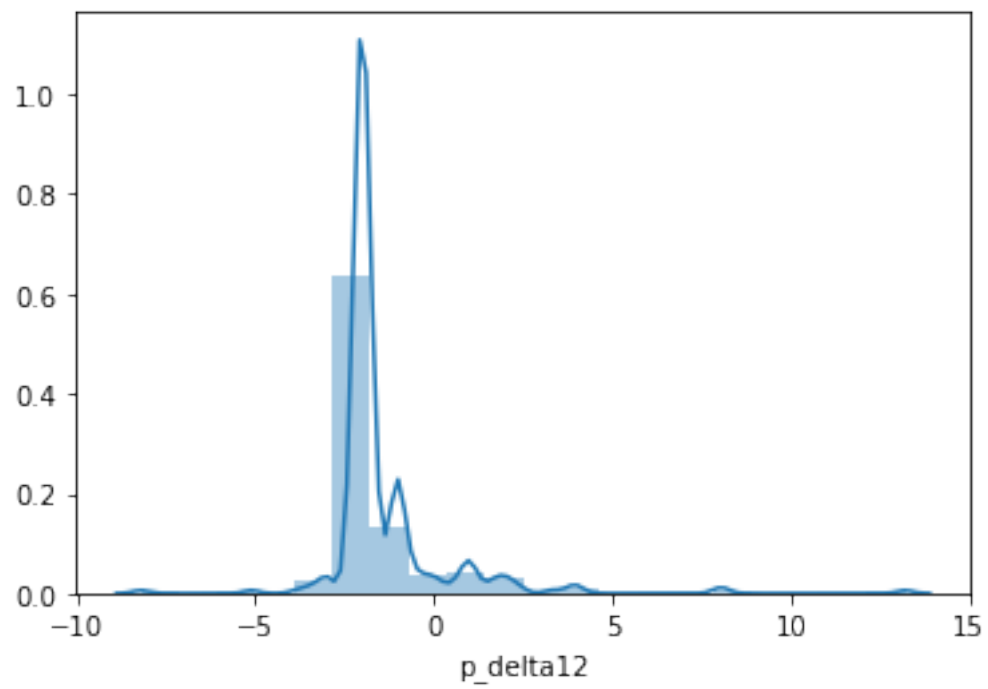
sns.distplot(df_meteo2['p_delta12'], bins = 20, kde = True)
plt.show()
```

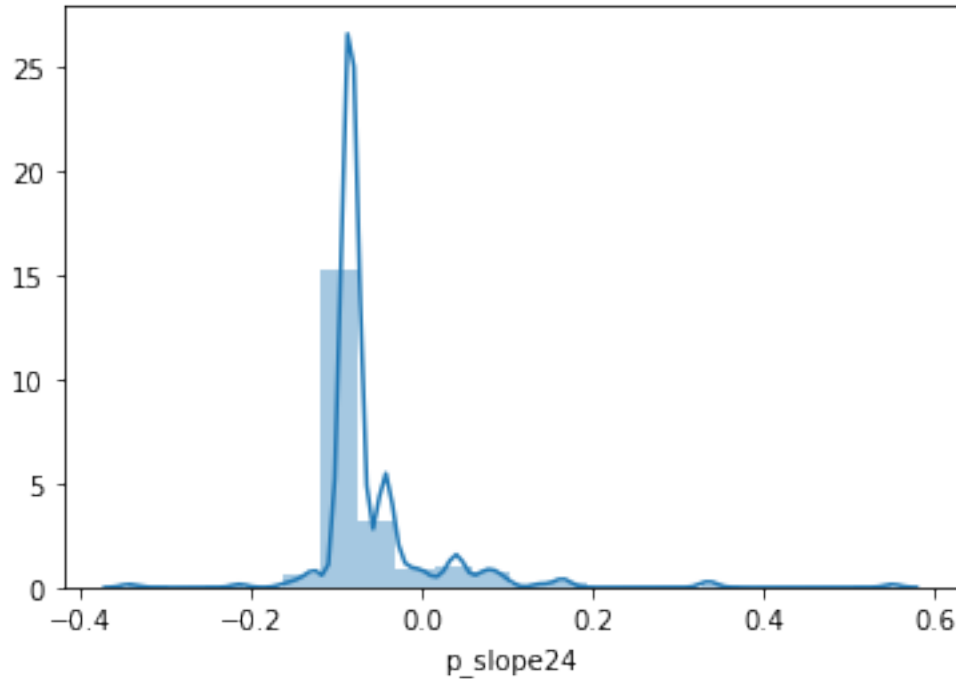
```
sns.distplot(df_meteo2['p_delta24'], bins = 20, kde = True)
plt.show()

sns.distplot(df_meteo2['p_slope24'], bins = 20, kde = True)
plt.show()

# sns.distplot(df_meteo2['t_delta12'], bins = 20, kde = True)
# plt.show()
```







### Creating Synthetic Negatives- Weighting Beachgoer Seasonality

To create synthetic negatives the same length as the existing dataset, the simplest approach is to use a random choice function pick non-attack occurrences that are not within the existing dataset. While this is easy to implement, there are some drawbacks to this. For this particular problem, the drawback is that shark attacks only occur when people are present. Furthermore, there are more or less people present at beaches in a) certain times of the year (seasonality) and b) in certain times of the day. If one was to randomly grab days using a random choice function, the randomization would not account for this seasonality and time of day where beachgoers are present.

So, how does one go about normalizing the random choice function to account for beachgoer seasonality? In this instance, data was taken from a publicly accessible dataset created by the University of West Florida in the form of monthly flight inflow data into Florida Coastal airports. Using this data, monthly weights were applied to the randomization to account for the seasonality of beachgoers.

It should be noted, that this weighting is not a perfect representation of when individuals go to the beach, but it does provide a directionally correct indication of beachgoer seasonality. Additionally, although this data is captured on a monthly basis, this particular project did not account for time-of-day (hourly) beach going patterns. This is a note that was created to be implemented in future iterations of this modeling to increase efficacy. Finally, in the case of temporal models, seasonality was not able to be implemented as the synthetic negatives were randomized on an hourly basis to match the dataset. The weather data was specifically where seasonally-adjusted synthetic negatives were implemented. The limitation of the temporal models' synthetic negatives is noted and will be fixed for future iterations.



```
[47]: df_2['month'] = pd.to_datetime(df_2['Day of Year']).dt.month
```

```
weight_m = df_2.groupby('month')['Airline Passengers'].sum().  
    ↪reindex(range(1,13), fill_value =0)  
weight_norm = weight_m/weight_m.sum()  
  
print(weight_norm, weight_m)
```

```
month  
1      0.060234  
2      0.062019  
3      0.083680  
4      0.086371  
5      0.091415  
6      0.103811  
7      0.104478  
8      0.094826  
9      0.079498  
10     0.087067  
11     0.075593  
12     0.071008  
Name: Airline Passengers, dtype: float64 month  
1      3909262  
2      4025125  
3      5430959  
4      5605609  
5      5932964  
6      6737491  
7      6780733  
8      6154342  
9      5159502  
10     5650741  
11     4906060  
12     4608478  
Name: Airline Passengers, dtype: int64
```

After creating my weights above, I now created my synthetic negatives based on those weightings. This allows me to skew the negative occurrences (my no-attack values) toward the summer months. Here, after creating these negative occurrences, I am joining them to my positive data set. I should note, also, that I am filling the negative occurrences with random weather values from the day I am choosing as well, so that we can compare like features in the positive set.

```
[65]: m= weight_norm.index  
w = weight_norm.values  
  
cols = ['date_at_location', 'county_match', 'temp', 'pres', 't_delta12',  
    ↪ 't_delta24', 'p_delta12', 'p_delta24',  
        't_slope12', 't_slope24', 'p_slope12', 'p_slope24', 't_nan', 'p_nan']
```

```

#df_meteo2 = df_meteo2.copy()
df_meteo2['month'] = pd.to_datetime(df_meteo2['date_at_location']).dt.month

df_pos2 = df_meteo2.copy()
df_pos2['attack_binary'] = 1

neg_m = np.random.choice(m, size = len(df_meteo2), p = w)

neg_row = []
for mm in neg_m:
    p = df_meteo2[df_meteo2['month'] == mm]
    if p.empty:
        p = df_meteo2
    r = p.sample(1, replace = True).iloc[0]
    neg_row.append(r[cols + ['month']])

df_neg2 = pd.DataFrame(neg_row).reset_index(drop = True)
df_neg2['attack_binary'] = 0

df_metmodel = pd.concat([df_pos2, df_neg2], ignore_index = True)
print(df_metmodel['attack_binary'].value_counts())
print(df_metmodel.head())

```

```

1    284
0    284
Name: attack_binary, dtype: int64

```

	attack_classification	outcome	county	date_at_location	Year	\
0	unprovoked attack	non-fatal	volusia county	2000-06-02	2000.0	
1	unprovoked attack	non-fatal	volusia county	2000-06-19	2000.0	
2	unprovoked attack	non-fatal	volusia county	2000-07-02	2000.0	
3	unprovoked attack	non-fatal	volusia county	2000-07-02	2000.0	
4	unprovoked attack	non-fatal	volusia county	2000-07-04	2000.0	

	time_of_attack	victim_activity	sunlight_conditions	\
0	14:55	body surfing, planing on waves	daylight	
1	08:30	wading	daylight	
2	17:02	standing still on bottom	daylight	
3	14:50	standing still on bottom	daylight	
4	18:00	wading	daylight	

	gen_weather	air_temp	...	p_delta12	p_delta24	\
0	clear	32.0	...	-2.0	-2.0	
1	data insufficient for judgement	32.0	...	-2.0	-2.0	
2	clear/calm	32.0	...	-2.0	-2.0	
3	data insufficient for judgement	32.0	...	-2.0	-2.0	
4	clear/calm	32.0	...	-2.0	-2.0	

	t_slope12	t_slope24	p_slope12	p_slope24	county_match	p_nan	t_nan	\
0	0.5	0.0	-0.166667	-0.083333	volusia county	1	1	
1	0.5	0.0	-0.166667	-0.083333	volusia county	1	1	
2	0.5	0.0	-0.166667	-0.083333	volusia county	1	1	
3	0.5	0.0	-0.166667	-0.083333	volusia county	1	1	
4	0.5	0.0	-0.166667	-0.083333	volusia county	1	1	

	attack_binary
0	1
1	1
2	1
3	1
4	1

[5 rows x 53 columns]

Here I am starting my EDA with the newly added synthetic negatives.

As you can see below, it seems a though there is a slight skew toward higher temperatures for attack occurences, as well as a slight skew toward changes in pressure over a 24 hour span, as well as the slope in the change of pressure over the same 24 hour span.

This is an interesting finding to me, and I wonder if we can actually prove that these are strong indicators of an attack vs non-attack in our models. So I will take note of these features for now and include them in my modeling.

```
[67]: plt.figure()
sns.boxplot(x = 'attack_binary', y = 'temp', data = df_metmodel)
plt.xticks([0,1], ['No-Attack', 'Attack'])
plt.title('Temp Dist on Attack vs No-Attack')
plt.show()

plt.figure()
sns.boxplot(x = 'attack_binary', y = 'pres', data = df_metmodel)
plt.xticks([0,1], ['No-Attack', 'Attack'])
plt.title('Pressure Dist on Attack vs No-Attack')
plt.show()

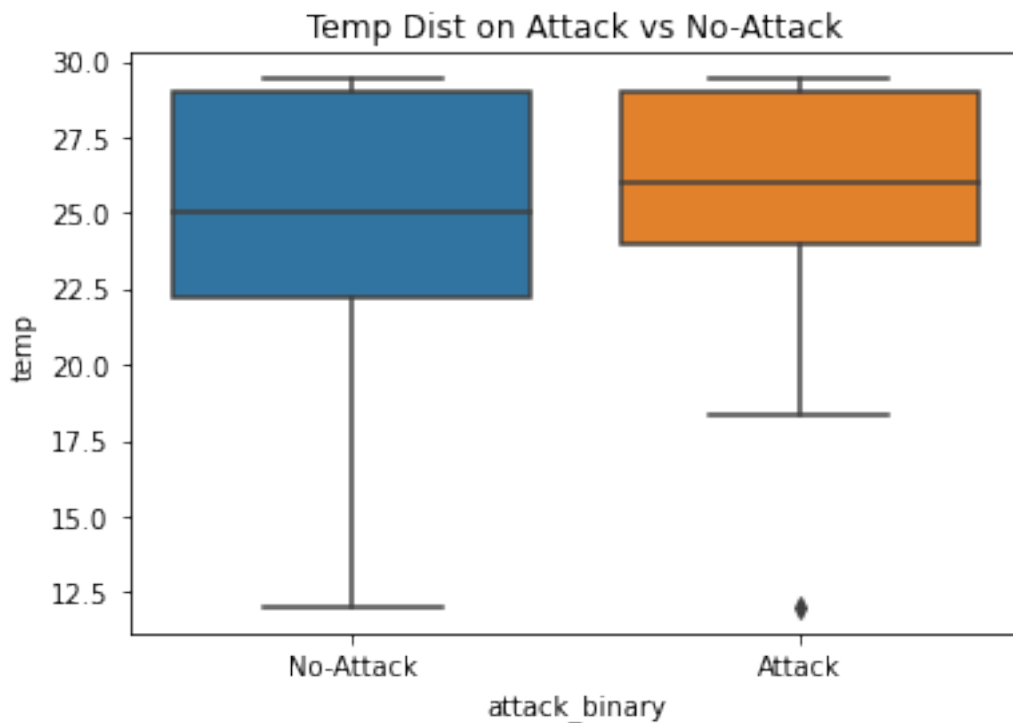
plt.figure()
sns.boxplot(x = 'attack_binary', y = 't_delta24', data = df_metmodel)
plt.xticks([0,1], ['No-Attack', 'Attack'])
plt.title('Temp 24 hr delta on Attack vs No-Attack')
plt.show()

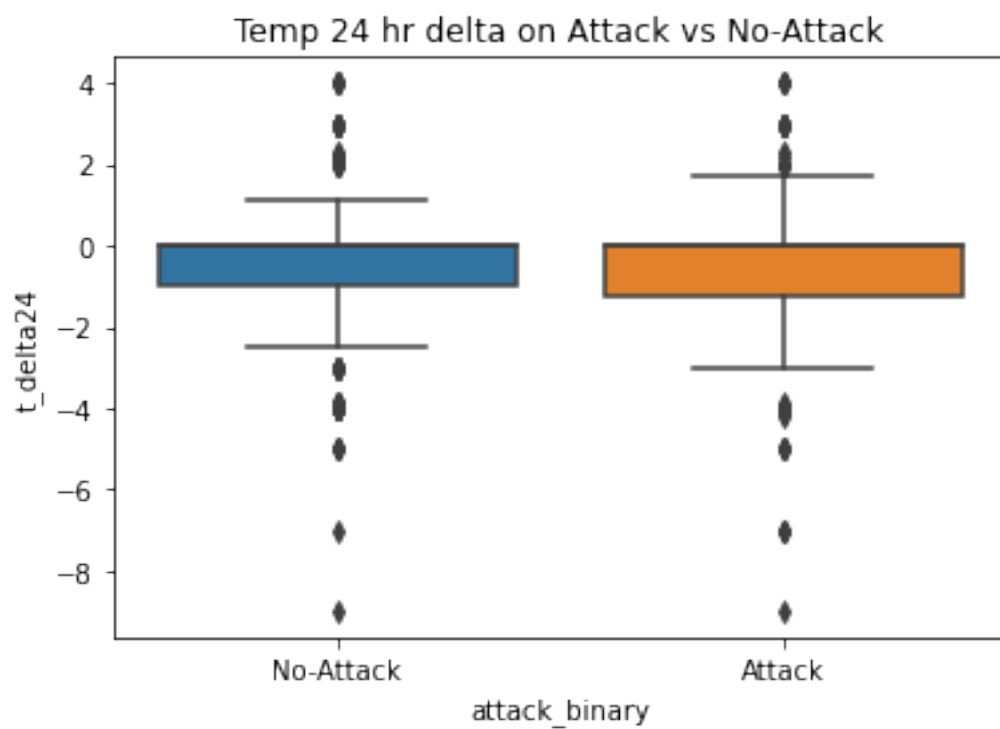
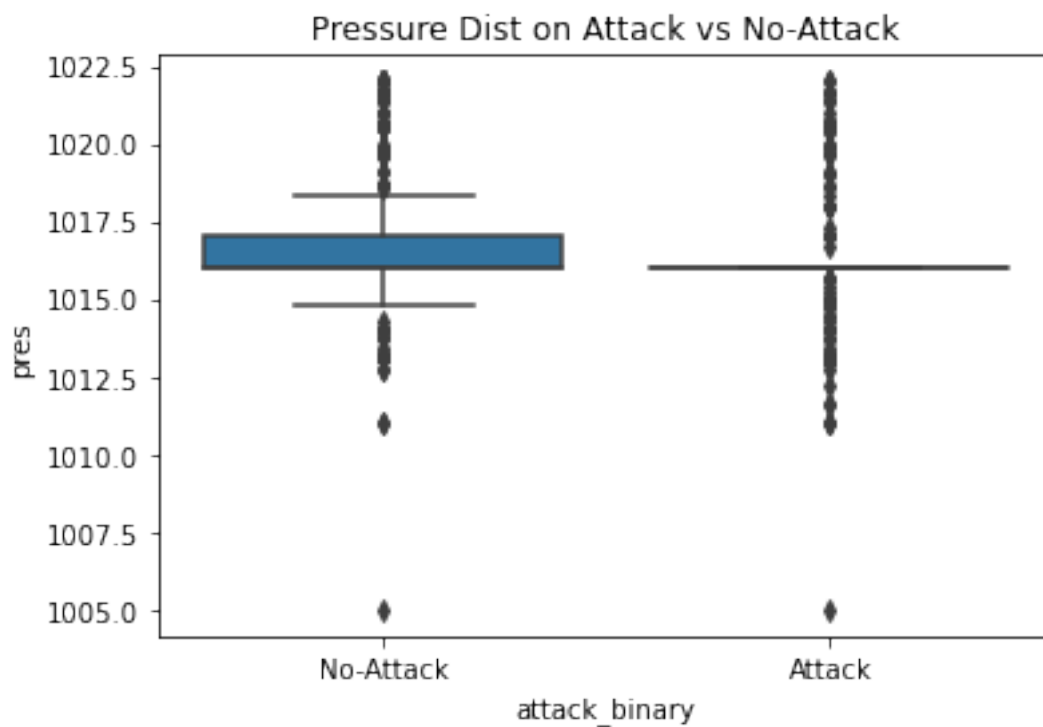
plt.figure()
sns.boxplot(x = 'attack_binary', y = 'p_delta24', data = df_metmodel)
plt.xticks([0,1], ['No-Attack', 'Attack'])
plt.title('Pressure 24 hr delta on Attack vs No-Attack')
```

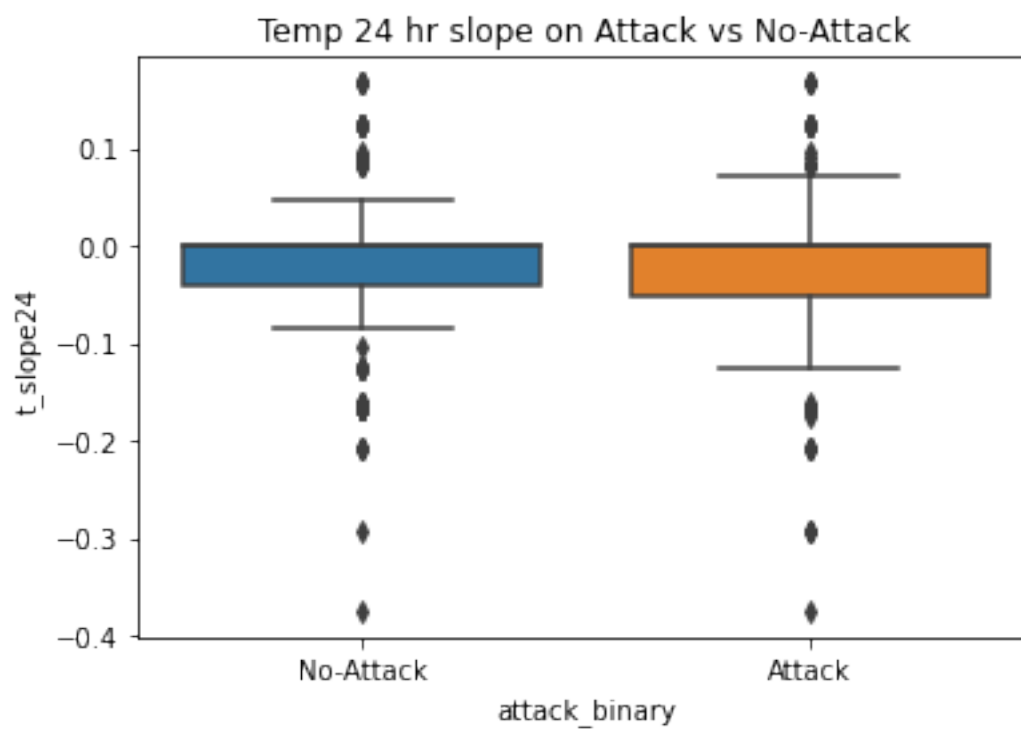
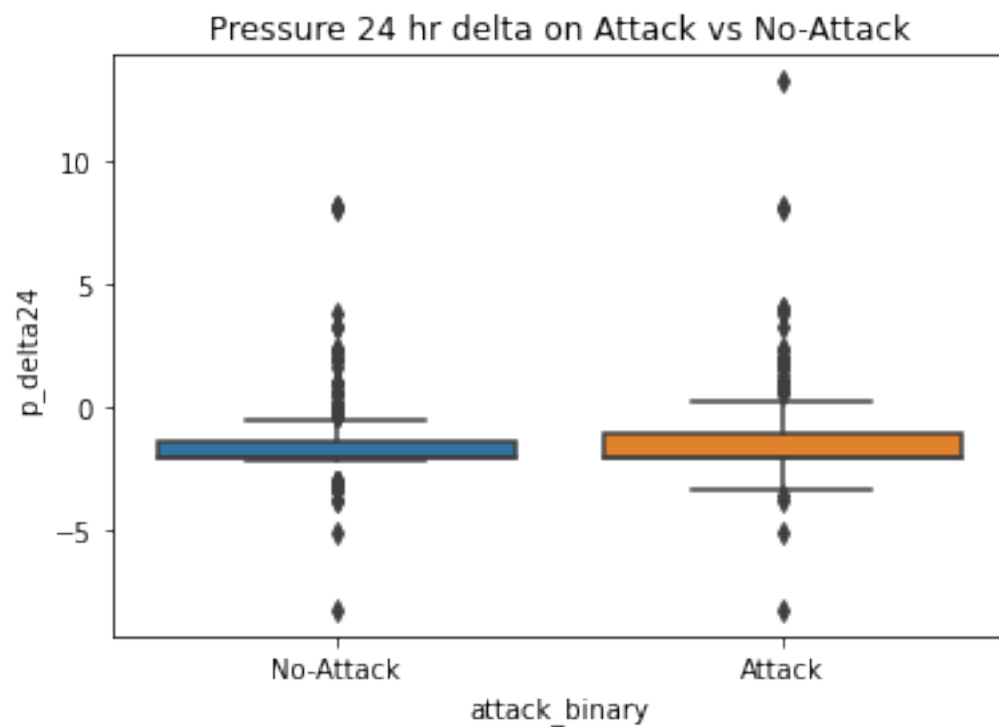
```
plt.show()

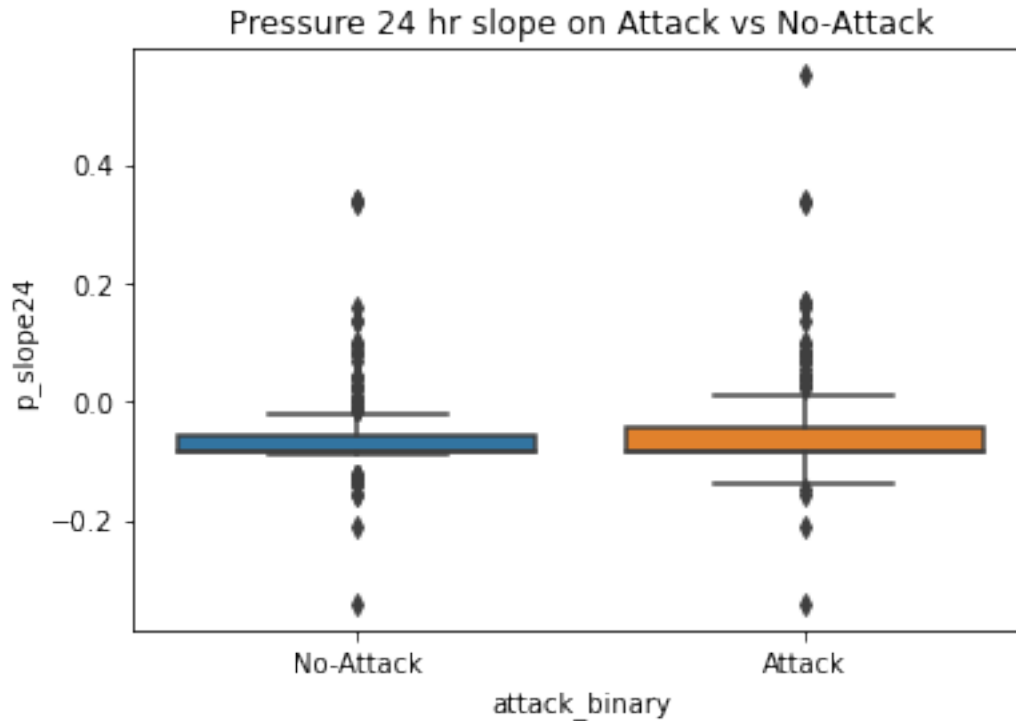
plt.figure()
sns.boxplot(x = 'attack_binary', y = 't_slope24', data = df_metmodel)
plt.xticks([0,1], ['No-Attack', 'Attack'])
plt.title('Temp 24 hr slope on Attack vs No-Attack')
plt.show()

plt.figure()
sns.boxplot(x = 'attack_binary', y = 'p_slope24', data = df_metmodel)
plt.xticks([0,1], ['No-Attack', 'Attack'])
plt.title('Pressure 24 hr slope on Attack vs No-Attack')
plt.show()
```









Again, another set of visualizations, this time limited to the three features I have honed in on: temperature, change in pressure over the past 24 hours, as well as the slope of this change. I think we see some variability here between these two features that we can investigate in our models

```
[68]: plt.figure()
sns.kdeplot(df_metmodel.loc[df_metmodel['attack_binary']==1, 'temp'], label = 'attack', shade = True)
sns.kdeplot(df_metmodel.loc[df_metmodel['attack_binary']==0, 'temp'], label = 'no attack', shade = True)
plt.title('Temp density of attack vs no attack')
plt.legend()
plt.show()

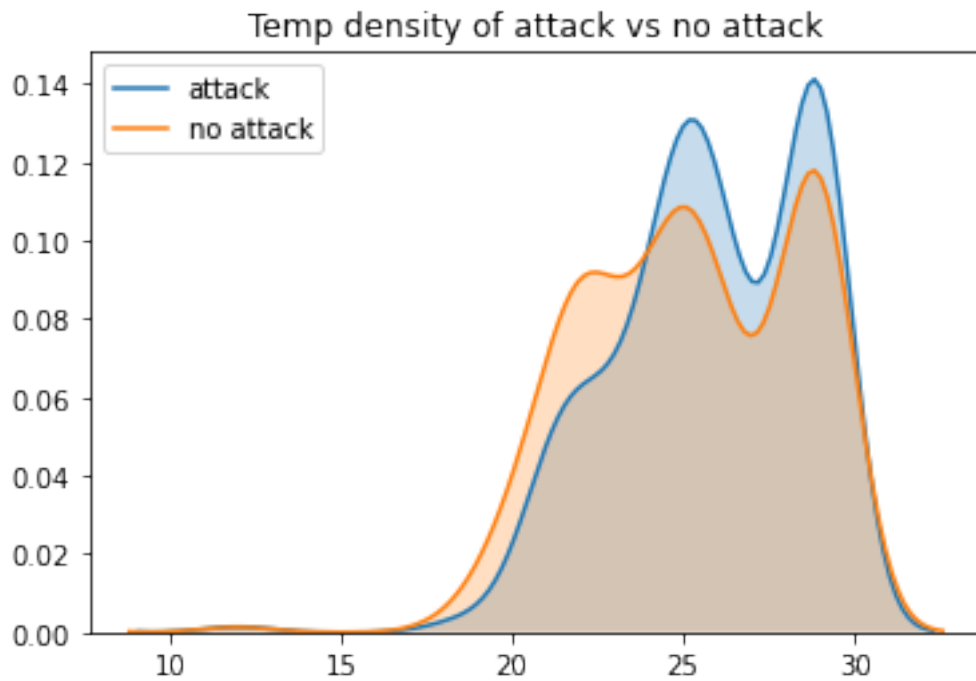
plt.figure()
sns.kdeplot(df_metmodel.loc[df_metmodel['attack_binary']==1, 'p_delta24'], label = 'attack', shade = True)
sns.kdeplot(df_metmodel.loc[df_metmodel['attack_binary']==0, 'p_delta24'], label = 'no attack', shade = True)
plt.title('Pressure 24 hr delta of attack vs no attack')
plt.legend()
plt.show()

plt.figure()
```

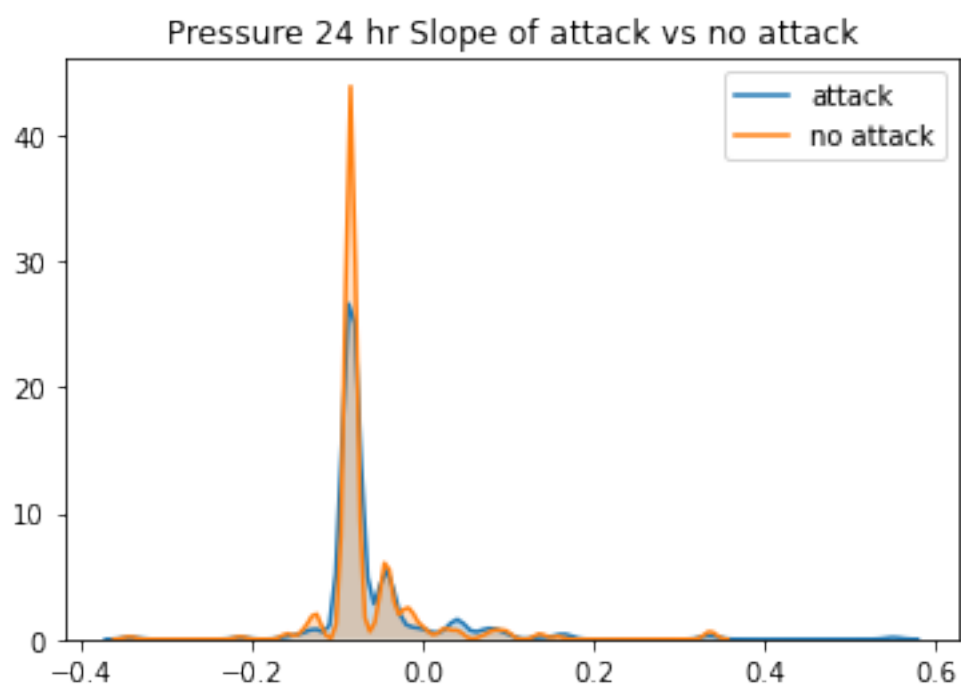
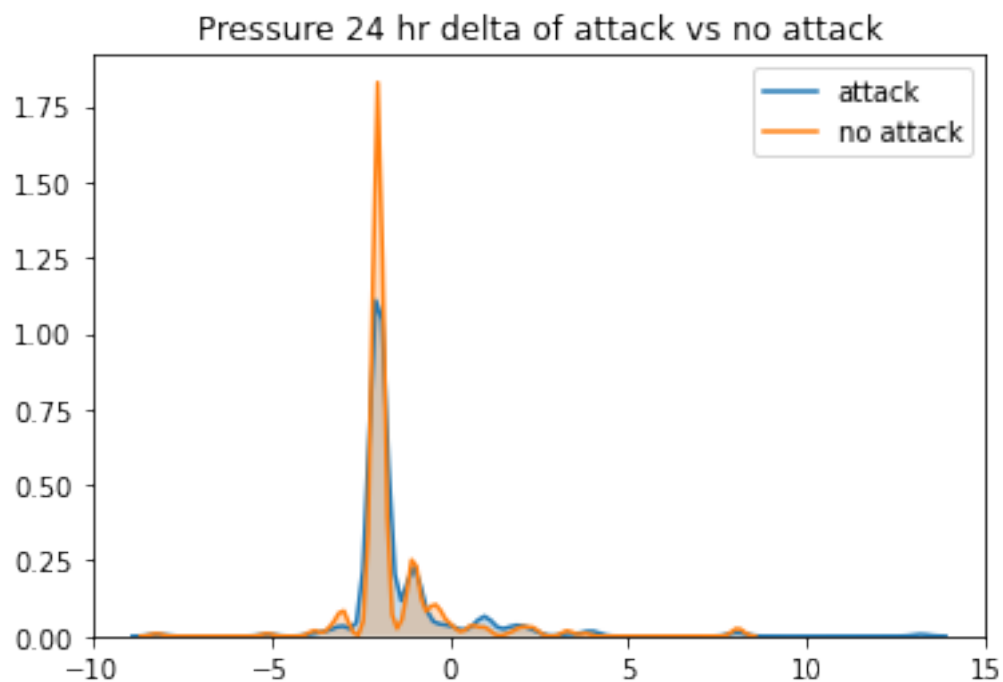
```

sns.kdeplot(df_metmodel.loc[df_metmodel['attack_binary']==1, 'p_slope24'],
            label = 'attack', shade = True)
sns.kdeplot(df_metmodel.loc[df_metmodel['attack_binary']==0, 'p_slope24'],
            label = 'no attack', shade = True)
plt.title('Pressure 24 hr Slope of attack vs no attack')
plt.legend()
plt.show()

```







### 5.3 Modeling Temperature, 24 Hour Pressure Change, and 24 Hour Pressure Slope

As you can see, my temperature turned out to be a poor indicator. My model really only predicted the outcome based on this feature little more than a coin flip could have. Lets move on to pressure change

```
[69]: #X_array = ['temp', 'p_delta24', 'p_slope24']

X = df_metmodel[['temp']]
y = df_metmodel['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.3,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(model.coef_)
```

	precision	recall	f1-score	support
0	0.54	0.56	0.55	86
1	0.54	0.52	0.53	85
accuracy			0.54	171
macro avg	0.54	0.54	0.54	171
weighted avg	0.54	0.54	0.54	171

```
[[48 38]
 [41 44]]
[[0.07423246]]
```

Pressure change, unfortunately turned out to be a poor indicator as well, again only being slightly better than an coin flip

```
[70]: X = df_metmodel[['p_delta24']]
y = df_metmodel['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.3,
↳stratify = y, random_state = 42)
```

```

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(model.coef_)

```

	precision	recall	f1-score	support
0	0.52	0.77	0.62	86
1	0.55	0.28	0.37	85
accuracy			0.53	171
macro avg	0.53	0.52	0.50	171
weighted avg	0.53	0.53	0.50	171

```

[[66 20]
 [61 24]]
[[0.06484879]]

```

As one would expect, the slope (similar to pressure change), performed poorly as well.

```

[71]: X = df_metmodel[['p_slope24']]
      y = df_metmodel['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y , test_size = 0.3,
↳stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(model.coef_)

```

	precision	recall	f1-score	support
0	0.52	0.77	0.62	86
1	0.55	0.28	0.37	85

accuracy			0.53	171
macro avg	0.53	0.52	0.50	171
weighted avg	0.53	0.53	0.50	171

```
[[66 20]
 [61 24]]
[[0.51984869]]
```

Finally, let's try combining all three features to see if we can create a stronger model. Unfortunately again, it looks like we are slightly better than a coin flip on this model, so I think we can conclude these features are not strong indicators of shark attack occurrence.

```
[72]: # from sklearn.preprocessing import StandardScaler
# from sklearn.pipeline import make_pipeline

X = df_metmodel[['temp', 'p_delta24', 'p_slope24']]
y = df_metmodel['attack_binary']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
→stratify = y, random_state = 42)

model = LogisticRegression(max_iter = 1000, class_weight = 'balanced')
#model = make_pipeline(StandardScaler(), LogisticRegression(max_iter = 1000,
→class_weight = 'balanced'))
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[: , 1]

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print(model.coef_)
print(roc_auc_score(y_test, y_prob))

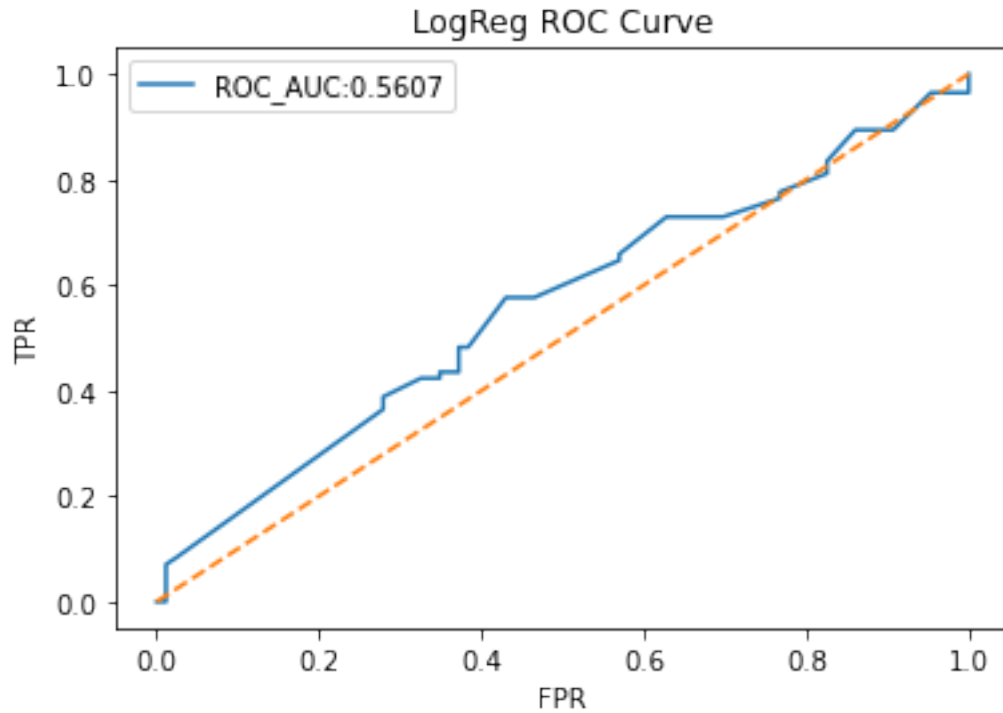
fpr, tpr, thresholds = roc_curve(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label = f"ROC_AUC:{roc_auc_score(y_test, y_prob):.4f}")
plt.plot([0,1], [0,1], linestyle = "--")
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LogReg ROC Curve')
plt.legend()
plt.show()
```

precision	recall	f1-score	support
-----------	--------	----------	---------

	0	0.55	0.63	0.59	86
	1	0.56	0.48	0.52	85
accuracy				0.56	171
macro avg		0.56	0.56	0.55	171
weighted avg		0.56	0.56	0.55	171

```
[[54 32]
 [44 41]]
[[0.08335977 0.09148068 0.0038117 ]]
0.5606703146374828
```



## 6 5. Results, Conclusions, Notes for Further Investigation, and Potential Applications of Findings

### Modeling Results

In this project, 10 separate Logistic regression models were created. They are as follows:

1. Shark Attack Occurrence ~ Lunar Illumination,
2. Shark Attack Occurrence ~ Month of Attack,
3. Shark Attack Occurrence ~ Hour of Attack,
4. Shark Attack Occurrence ~ Hourly Bucket (Time of Day),

5. Shark Attack Occurrence ~ Lunar Phase,
6. Shark Attack Occurrence ~ Month of Attack + Hourly Bucket (Time of Day),
7. Shark Attack Occurrence ~ Temperature,
8. Shark Attack Occurrence ~ 24 Hour Pressure Change,
9. Shark Attack Occurrence ~ 24 Hour Pressure Rate of Change,
10. Shark Attack Occurrence ~ Temp + 24 Hour Pressure Change + 24 Hour Pressure Rate of Change

### Model Performance Summary and Comparison

Model	Encoding	ROC-AUC	Notes
<b>Shark Attack Occurrence ~ Lunar Illumination</b>	Continuous	~0.50	Weak; nearly random performance.
<b>Shark Attack Occurrence ~ Month of Attack</b>	Cyclical (sin/cos)	~0.69	Not a strong predictor; small seasonal effect visible.
<b>Shark Attack Occurrence ~ Hour of Attack</b>	Cyclical (sin/cos)	~0.76	Moderate strength; improves over month alone.
<b>Shark Attack Occurrence ~ Hourly Bucket (Time of Day)</b>	Categorical	~0.78	Strongest single predictor; captures diurnal patterns well.
<b>Shark Attack Occurrence ~ Lunar Phase</b>	Categorical (New, Full, etc.)	~0.45	Weak; nearly random performance.
<b>Shark Attack Occurrence ~ Month of Attack + Hourly Bucket</b>	One-hot month + Categorical	~0.80	Best performing combined model; strong diurnal + seasonal signal.
<b>Shark Attack Occurrence ~ Temperature</b>	Continuous (°C)	~0.53	Weak; nearly random performance.
<b>Shark Attack Occurrence ~ 24 Hour Pressure Change</b>	Continuous ( $\Delta$ pressure, hPa)	~0.54	Weak; nearly random performance.
<b>Shark Attack Occurrence ~ 24 Hour Pressure Rate of Change</b>	Continuous (slope, hPa/hr)	~0.54	Weak; nearly random performance.
<b>Shark Attack Occurrence ~ Temp + Pressure Features</b>	Continuous (3 features)	~0.56	Adding multiple meteorological features doesn't improve model.

For the sake of brevity, only two of the most significant results will be listed in this section from both the Temporal and Weather category, but all results are in the above notebook.

For the temporal data “Shark Attack Occurrence ~ Month of Attack + Hourly Bucket (Time of Day)” yielded the most significant results. As you can see by the results, combining these two features yielded a combined ROC AUC score of ~80%. This is a significant outcome, and means this combo model has a strong chance of predicting shark attack occurrence vs non-occurrence.

In the predicted probability charts for the monthly and hourly data in this model, one can see that the distribution of attack probability is similar to the real-world distribution of shark attacks that

was seen in the EDA visualizations earlier. This is a robust result, because it shows the model is predicting in a similar fashion to real world outcomes.

Additionally, providing the hourly and monthly distribution of attack vs non-attack below to visually see how the model parsed through positive and negative events in the test and train data:

For the Weather data, “Shark Attack Occurrence  $\sim$  Temp + 24 Hour Pressure Change + 24 Hour Pressure Rate of Change” yielded the most significant result. Unfortunately, this model was little better than a coin flip in predicting shark attack occurrences from these three variables stacked, with an ROC AUC score barely over 50%

## Discussion of Results

Looking at the results of the modeling done on this dataset, it is easy for one to tell that the temporal indicators provided the results of the highest value. Looking at the overall ROC AUC score of  $\sim$ 80% on the Hourly and Month day combination, this model provides high efficacy and will accurately predict the occurrence of an attack 76% of the time (accuracy score) when given date-time info. In terms of which model should be iterated upon “Shark Attack Occurrence  $\sim$  Month of Attack + Hour of Attack” is the most promising.

Other temporal data features looked promising as well, but may need further refinement. Lunar phases and lunar illumination, which were previously hypothesized to be a potential indicator, but yielded poor results when modelled. This information in itself is useful though in that it tells us that Lunar Phases may have no correlation whatsoever with shark attack occurrences. In this respect, Lunar temporal data requires further investigation.

In regards to the meteorological features- Those that were chosen for modeling (Temperature, 24 Hour Pressure Delta, and 24 Hour Pressure Rate of Change) were extremely poor indicators of shark attack occurrence. There may be further credence to the use of these features as indicators, if modeling is enhanced, but for the time being they lack a significance in the current iteration.

## Limitations of Current Model

One notable limitation for the current model is that synthetic negative events for temporal data were completely randomly generated and in future iterations, implementing seasonality and time-of-day into the random choice function would increase model efficacy. Additionally time-of-day generation in both random choice functions for temporal and weather data, if beachgoer hourly data is captured would increase model efficacy.

Furthermore, lack of expertise in ichthyology and shark research could have caused disregard for potential in other features that could increase the model efficacy. In future iterations, working with the ISAF researchers and gaining first-hand knowledge in shark behavior could help direct additional features to include.

Finally, bias in the geographical nature of data, being that this subset only included Florida Atlantic coast attacks could have very well skewed the results of certain models. Running a global or continent based model could drown out these regional biases. Additionally splitting the data along these regional lines could yield interesting and different results.

## Future Applications and Model Iterations

One potential application of a model such as this, when further refined to include additional features to increase overall efficacy, would be the use of a risk indicator of sorts for shark attack occurrence given a specific set of circumstances. Think of a “Smokey the Bear: Wildfire Risk”-type indicator

that could provide the generalized risk of a shark attack occurrence. This model could be further refined to add features such as geographical location, tidal and current affects, storm systems in nearby area, species indicators, and so on. With further expert input on the model features to be included, one could more greatly refine to provide individuals with valuable risk information when making decisions about their beach going activities.

## 7 References

Florida Museum. (n.d.). Beach injuries and fatalities: Comparing risks. Florida Museum of Natural History. Retrieved September 25, 2025, from <https://www.floridamuseum.ufl.edu/shark-attacks/odds/compare-risk/beach-injuries-fatalities/>

Florida Museum. (n.d.). International Shark Attack File. Florida Museum of Natural History. Retrieved September 25, 2025, from <https://www.floridamuseum.ufl.edu/shark-attacks/>

University of West Florida, Haas Center. (n.d.). Tourism indicators. University of West Florida. Retrieved September 25, 2025, from <https://uwf.edu/centers/haas-center/explore-the-economy/tourism-indicators/>

French, L. A., Midway, S. R., Evans, D. H., & Burgess, G. H. (2021). Shark side of the moon: Are shark attacks related to lunar phase? *Frontiers in Marine Science*, 8, 745221.<https://doi.org/10.3389/fmars.2021.745221>

OpenAI. (2025). ChatGPT (version GPT-5) [Large language model]. OpenAI. (Used for error trace debugging and function usage examples.)

Google DeepMind. (2025). Gemini [Large language model]. Google DeepMind. (Used for error trace debugging and function usage examples.)

Special acknowledgement: Dr. Gavin Naylor and Joseph Miguez, Florida Program for Shark Research, for providing ISAF data.

```
[2]: #df_clean.to_csv('df_clean_here', index = False)
```

```
[ ]:
```